

UNIX System Programming

Directories and File System

❖ Objectives

- look at how to program with directories
- briefly describe the UNIX file system

Overview

1. Directory Implementation
2. Subdirectory Creation
3. “.” and “..”
4. `mkdir()`
5. `rmdir()`
6. Reading Directories
7. `chdir()`
8. `getcwd()`
9. Links

1. Directory Implementation

❖ A UNIX directory is a *file*:

- it has an owner, group owner, size, access permissions, etc.
- many file operations can be used on directories

❖ Differences:

- modern UNIXs have special directory operations
e.g. `opendir()`, `readdir()`

Directory Structure

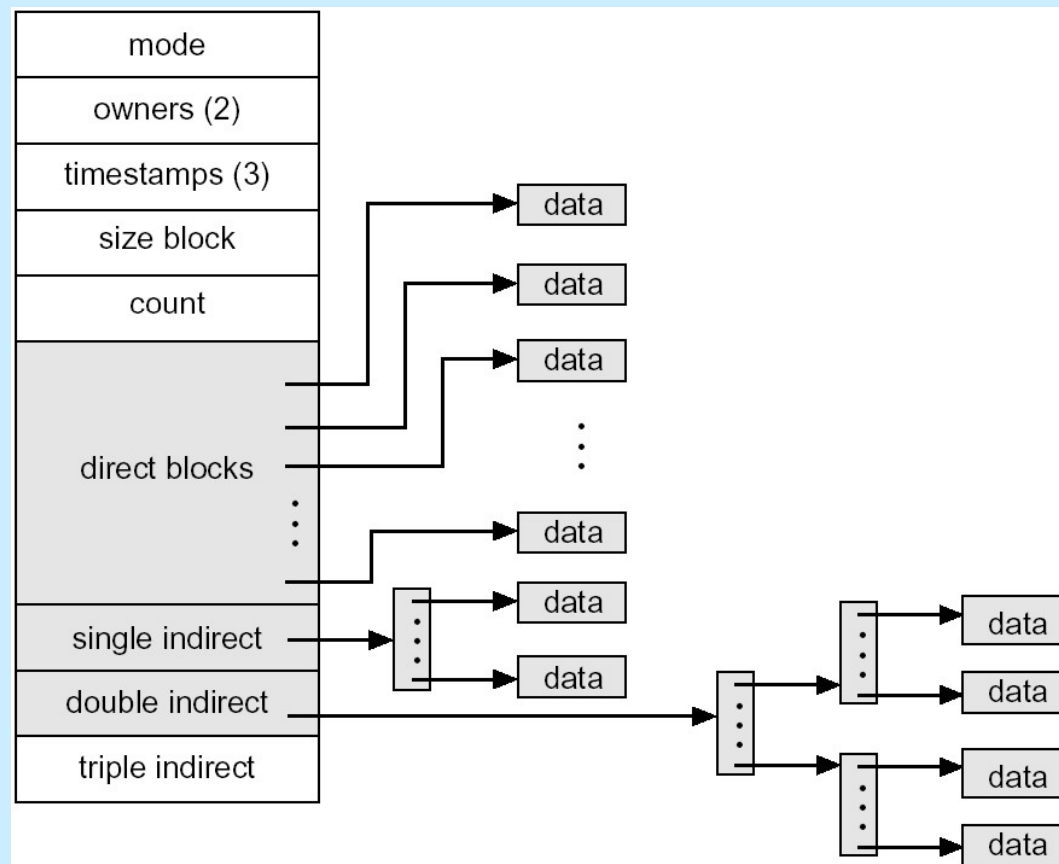
- ❖ A directory 'file' is a sequence of lines; each line holds an *i-node number* and a file name.
- ❖ The data is stored as binary, so we cannot simply use **cat** to view it

keeps track of what 'blocks' store your file's data

❖ *I-node*:

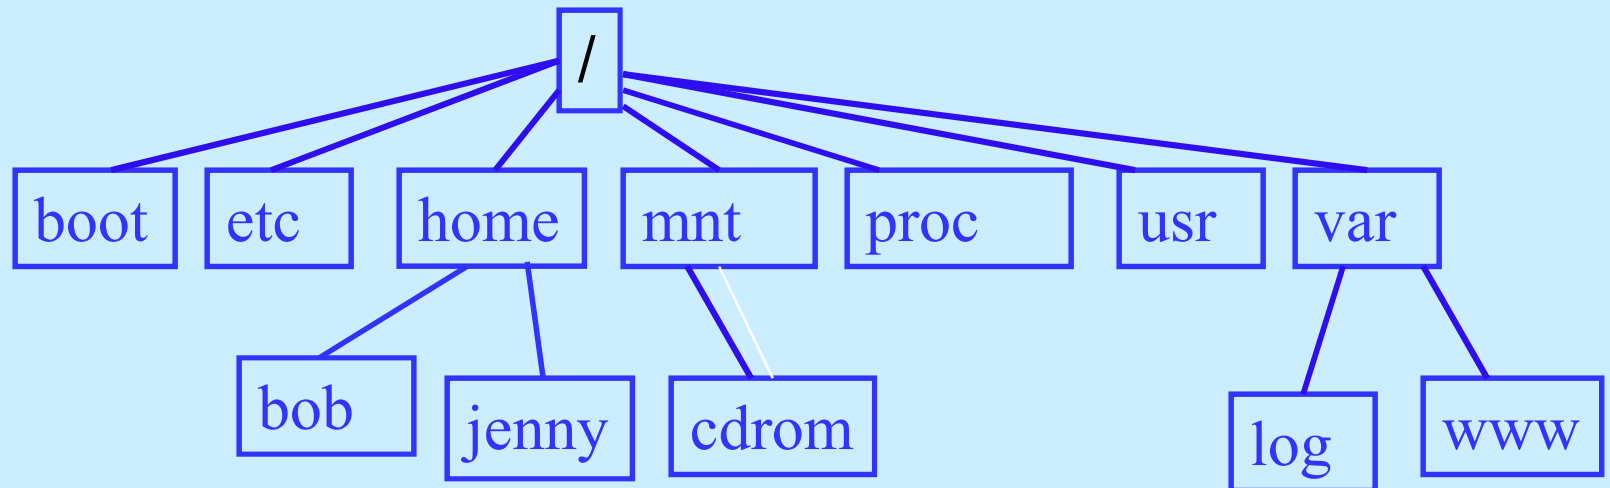
- The administrative information about a file is kept in a structure known as an *inode*.
 - ◆ Inodes in a file system, in general, are structured as an array known as an *inode table*.
- An inode number, which is an index to the inode table, *uniquely identifies* a file in a file system.

i-node and Data Blocks



UNIX directory hierarchy

- ❖ UNIX does not have the concept of separate drives as exist in the Windows world.
- ❖ All devices, when mounted, are accessed through a directory structure which starts at the / (root directory).



loki.trentu.ca

/ is the top of the tree

```
[jacques@loki ~]$ ls /
```

```
boot  etc  lib  lost+found  mnt  proc  run  srv  tmp  var
bin  dev  home  lib64  media  opt  root  sbin  sys  usr
```

```
[jacques@loki ~]$ df -h    df - disk subsystem
```

Filesystem	Size	Used	Avail	Use%	Mounted on
<u>/dev/sdb1</u>	15G	6.7G	7.2G	49%	/
devtmpfs	902M	0	902M	0%	/dev
tmpfs	916M	84K	916M	1%	/dev/shm
tmpfs	916M	97M	819M	11%	/run
tmpfs	916M	0	916M	0%	/sys/fs/cgroup
<u>/dev/sda1</u>	19G	301M	17G	2%	/home
tmpfs	184M	16K	184M	1%	/run/user/42
tmpfs	184M	0	184M	0%	/run/user/1000

Directory Creation

❖ “`mkdir uga`” causes:


- the creation of a `uga` directory file and an i-node for it
- an i-node number and name are added to the parent directory file

:	
120	“fred.html”
207	“abc”
135	“bookmark.c”
201	“ uga ”
:	

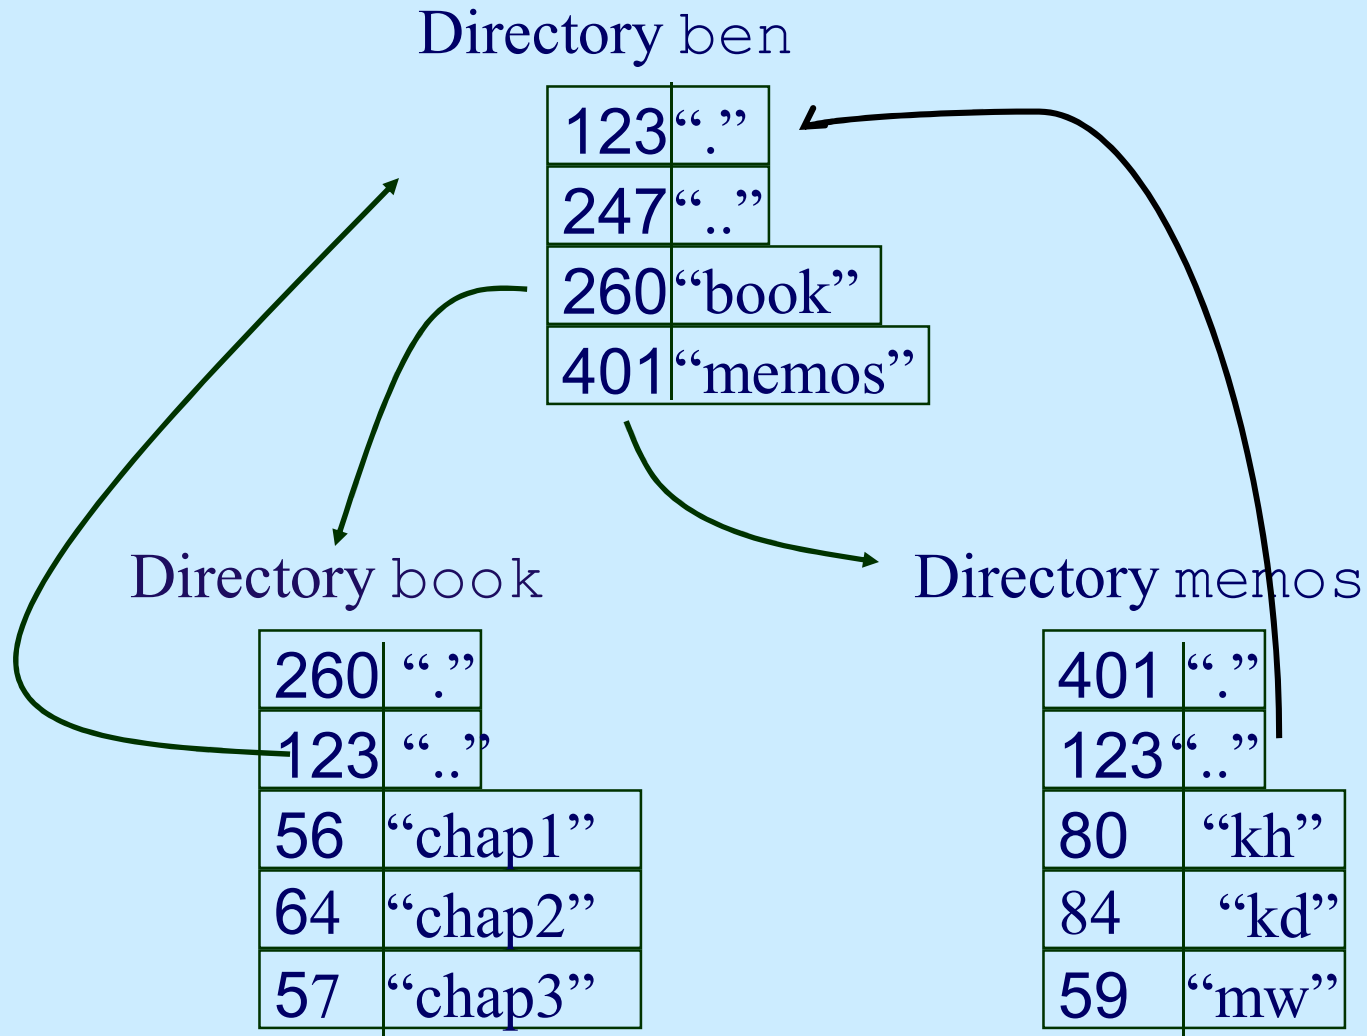
“.” and “..”

- ❖ In UNIX, the single period “.” represents the current directory.
- ❖ The double period “..” represents the parent directory.
- ❖ “.” and “..” are stored as ordinary file names with i-node numbers pointing to the correct directory files.

❖ Example:



In more detail:



mkdir()

- ❖ `#include <sys/types.h>`
`#include <fcntl.h>`
`#include <unistd.h>`

```
int mkdir(char *pathname, mode_t mode);
```

- ❖ Creates a new directory with the specified **mode**: return 0 if ok, -1 on error
- ❖ “.” and “..” entries are added automatically
- ❖ **mode** must include execute permissions so the user(s) can use **cd**.

e.g. 0755

read write execute
r w x
1 1 1 (binary)
= 7 decimal

rmkdir()

- ❖ `#include <unistd.h>`
`int rmkdir(char *pathname);`
- ❖ Delete an empty directory;
return 0 if ok, -1 on error.
- ❖ Will delay until other processes have stopped
using the directory.

Reading Directories

```
#include <sys/types.h>
#include <dirent.h>
```

returns a
pointer if ok,
NULL on error

```
DIR *opendir(char *pathname);
```

```
struct dirent *readdir(DIR *dp);
```

```
int closedir(DIR *dp);
```

returns a
pointer if ok,
NULL at end
or on error

dirent and DIR

❖ struct dirent

```
{
    long d_ino;                /* i-node number */
    char d_name[NAME_MAX+1];  /* fname */
    off_t d_off;               /* offset to next rec */
    unsigned short d_reclen;   /* record length */
}
```

❖ DIR is a directory stream (similar to FILE)

- when a directory is first opened, the stream points to the first entry in the directory

Example: listdir.c

List the contents of the current directory.

```
#include <stdio.h>
#include <dirent.h>

int main()
{
    DIR *dp;
    struct dirent *dir;

    if( (dp = opendir(".")) == NULL )
    {
        fprintf( stderr, "Cannot open dir\n" );
        exit(1);
    }

    /* read entries */
    while( (dir = readdir(dp)) != NULL )
    {
        /* ignore empty records */
        if( dir->d_ino != 0 )           if i node is not zero then print out the file name
            printf( "%s\n", dir->d_name );
    }

    closedir( dp );
    return 0;
} /* end main */
```


chdir()

```
#include <unistd.h>
int chdir( char *pathname );
int fchdir( int fd );
```

- ❖ Change the current working directory (*cwd*) of the calling process; return 0 if ok, -1 on error.

Example: cd to /tmp

❖ Part of `to_tmp.c`:

```
:  
if( chdir("/tmp" ) < 0 )  
    printf( "chdir error\n" ) ;  
else  
    printf( "In /tmp\n" ) ;
```

Directory Change is Local

❖ The directory change is limited to within the program.

❖ e.g.

```
$ pwd  
/usr/lib  
$ to_tmp  
In /tmp  
$ pwd  
/usr/lib
```

```
/* from last slide */
```

getcwd()

- ❖ `#include <unistd.h>`
`char *getcwd(char *buf, int size);`
- ❖ Store the *cwd* of the calling process in `buf`;
return `buf` if ok, `NULL` on error.
- ❖ `buf` must be big enough for the pathname string
(`size` specifies the length of `buf`).

Example

```
#include <stdio.h>
#include <unistd.h>
#include <dirent.h> /* for NAME_MAX */

int main()
{
    char name[NAME_MAX+1];

    if( getcwd( name, NAME_MAX+1 ) == NULL )
        printf( "getcwd error\n" );
    else
        printf( "cwd = %s\n", name );
    :
```

2. Links

- 2.1 What is a Link?
- 2.2 Creating a Link
- 2.3 Seeing Links
- 2.4 Removing a Link
- 2.5 Symbolic Links
- 2.6 Implementation

2.1. What is a Link?

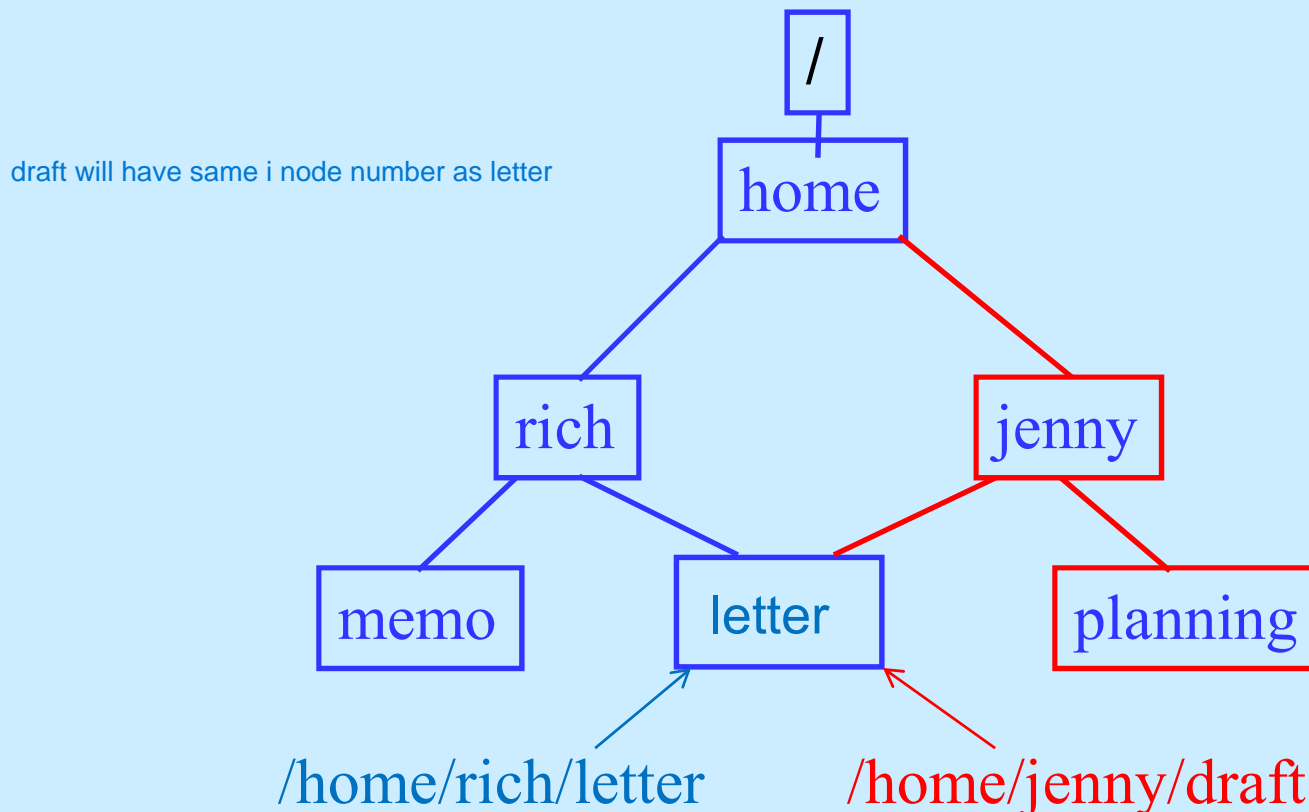
- ❖ A link is a pointer to a file.
- ❖ Useful for sharing files:
 - a file can be shared by giving each person their own link (pointer) to it.

2.2. Creating a Link

syntax: `ln existing-file new-pointer`

❖ Jenny types:

`ln draft /home/rich/letter`



❖ Changes to a file affects every link:

```
$ cat file_a  
This is file A.  
$ ln file_a file_b  
$ cat file_b  
This is file A.
```

```
$ vi file_b  
:
```

```
$ cat file_b  
This is file B after the change.  
$ cat file_a  
This is file B after the change.
```

2.3. Seeing Links

❖ Compare status information:

```
$ ls -l file_a file_b file_c file_d

-rw-r--r-- 2 dk1 33 May 24 10:52 file_a
-rw-r--r-- 2 dk1 33 May 24 10:52 file_b
-rw-r--r-- 1 dk1 16 May 24 10:55 file_c
-rw-r--r-- 1 dk1 33 May 24 10:57 file_d
```

❖ Look at inode number:

```
$ ls -li file_a file_b file_c file_d
```

-i to print i node numbers

```
3534 file_a    3534 file_b
5800 file_c    7328 file_d
```

2.4. Removing a Link

- ❖ Deleting a link does not remove the file.
- ❖ Only when the file *and* every link is gone will the file be removed.

2.5. Symbolic Links

- ❖ The links described so far are often called *hard links*
 - a hard link is a pointer to a file which must be on the *same* file system
- ❖ A *symbolic link* on separate physical hard drive is an *indirect pointer* to a file
 - it stores the pathname of the pointed-to file
 - it can link **across** file systems

Lets say we have a system with a number of different hard drives.

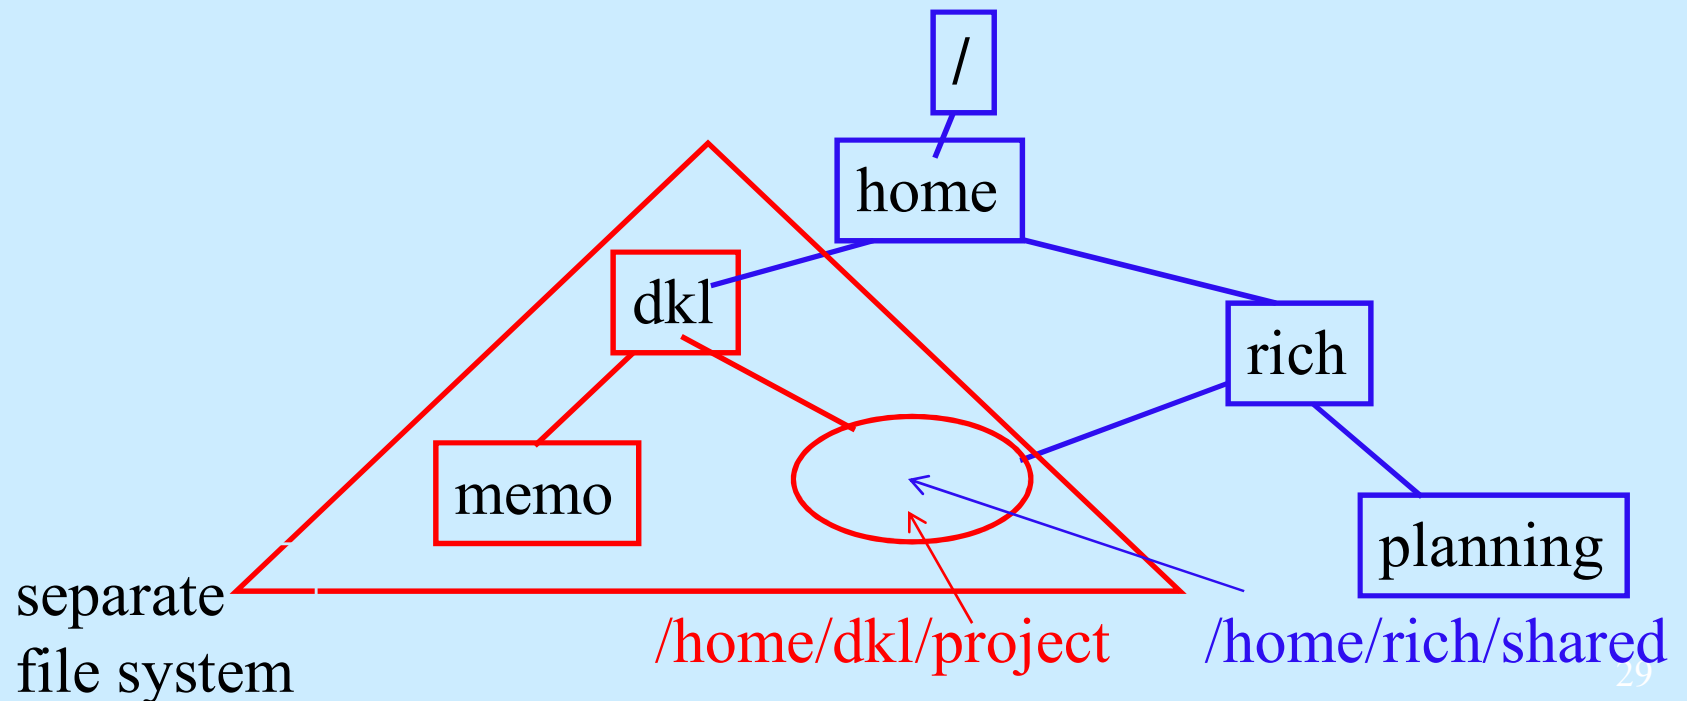
The **dkl** group has their own drive and have mounted off of the /home directory.

when Richard wants to access the file he has to type in /home/dkl/project every time.

He can create a link with a name in his own directory that he can access without having to specify a whole path.

❖ Richard types:

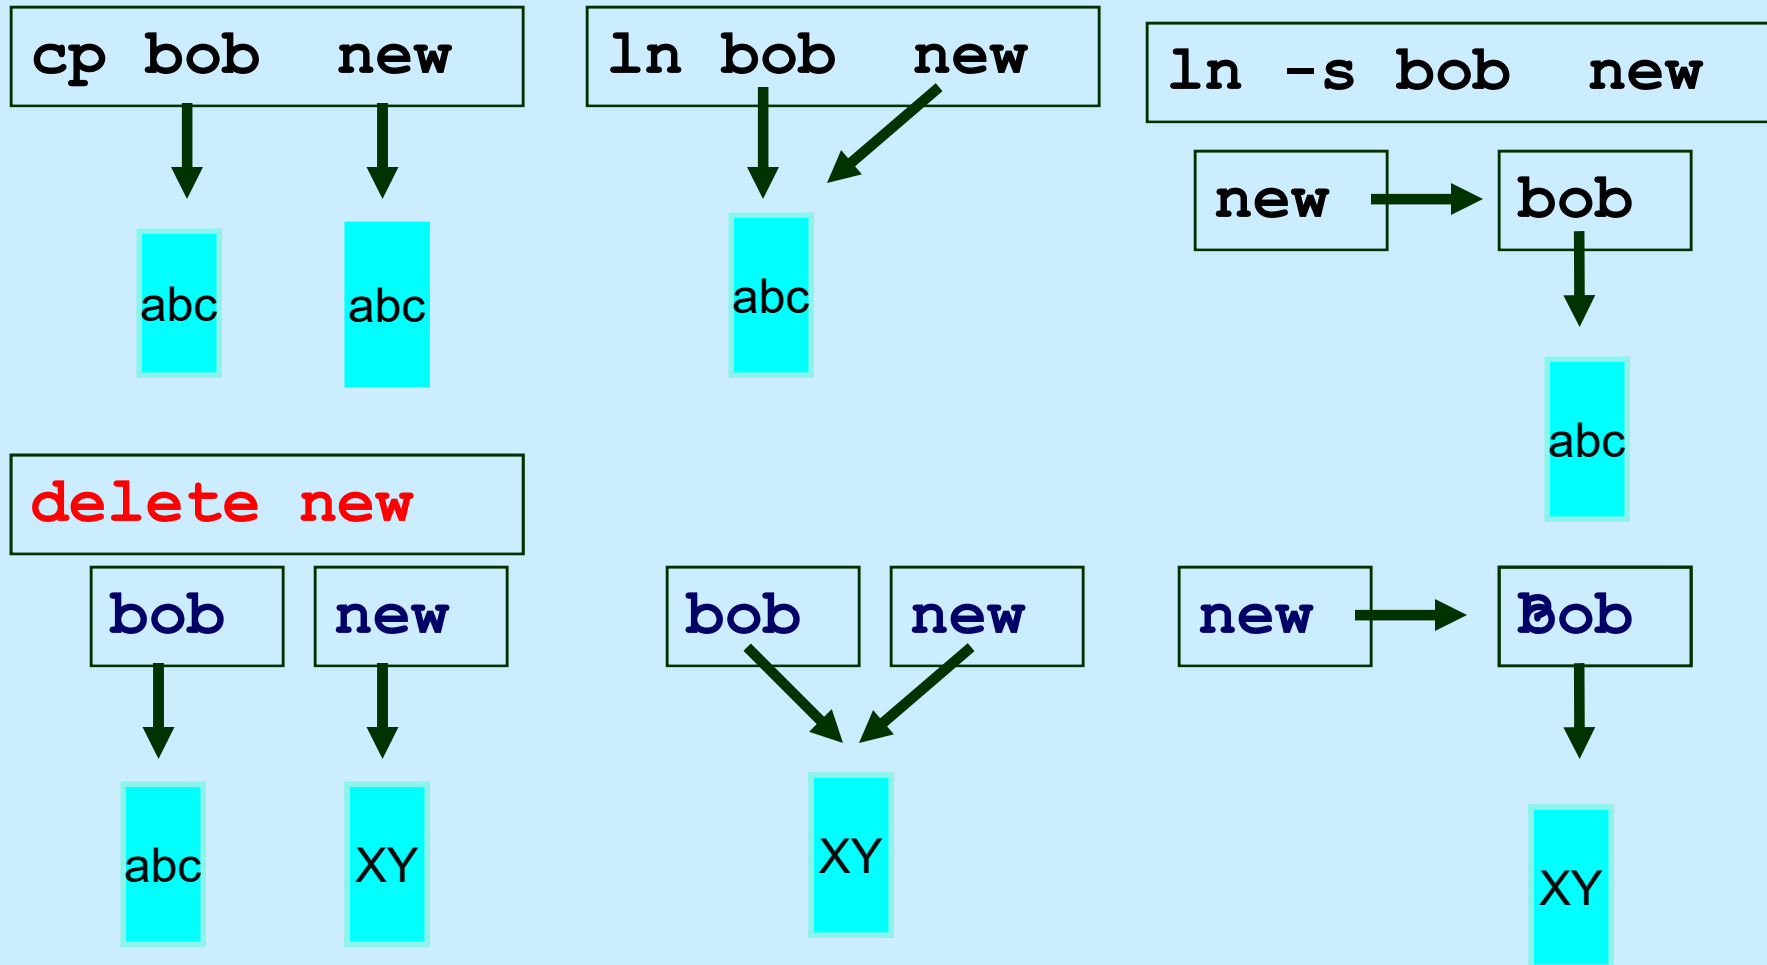
```
ln -s shared /home/dkl/project
```



❖ Symbolic links are listed differently:

```
$ ln -s pics /home/mh/img  
$ ls -lF pics /home/mh/img  
drw-r--r-- 1 dkl staff 981 May 24 10:55 pics  
lrwxrwxrwx 1 dkl staff  4 May 24 10:57/home/mh/img --> pics
```

2.6 Link Creation, Update & Removal



2.7 link() and unlink()

```
#include <unistd.h>
int link( const char *oldpath, const char *newpath );
```

❖ Meaning of:

```
link( "abc", "xyz" )
```

:

120	"fred.html"
207	"abc"
135	"bookmark.c"
207	"xyz"

:

- ❖ `unlink()` clears the **directory record**
 - usually means that the i-node number is set to 0
- ❖ The i-node is only deleted when the **last** link to it is removed; the data block for the file is also deleted (reclaimed) & no process have the file opened

Example: unlink

```
#include <stdio.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <fcntl.h>

int main(void)
{
    if( open( "tempfile", O_RDWR ) < 0 )
    {
        perror( "open error" );
        exit( 1 );
    }
    if( unlink( "tempfile" ) < 0 )
    {
        perror( "unlink error" );
        exit( 1 );
    }
    printf( "file unlinked\n" );
    exit(0);
}
```

symlink()

```
#include <unistd.h>
```

```
int symlink(const char *oldpath, const char *newpath) ;
```

- ❖ Creates a symbolic link named *newpath* which contains the string *oldpath*.
- ❖ Symbolic links are interpreted at run-time.
- ❖ Dangling link – may point to a non-existing file.
- ❖ If *newpath* exists it will not be overwritten.

readlink()

```
#include <unistd.h>
```

```
int readlink( const char *path, char *buf, size_t bufsiz);
```

❖ Read value of a **symbolic link** (does *not* follow the link).

- Places the contents of the symbolic link *path* in the buffer *buf*, which has size *bufsiz*.
- Does not append a NULL character to *buf*.

❖ Return value

- The count of characters placed in the buffer if it succeeds.
 - ◆ -1 if an error occurs.