


UNIX System Programming

Signals



Overview

1. Definition
2. Signal Types
3. Generating a Signal
4. Responding to a Signal
5. Common Uses of Signals
6. Timeout on a `read()`
7. POSIX Signal Functions
8. Interrupted System Calls
9. System Calls inside Handlers
10. More Information

1. Definition

- ❖ A signal is an *asynchronous* event which is delivered to a process.
- ❖ Asynchronous means that the event can occur at any time
 - may be unrelated to the execution of the process
 - e.g. user types `ctrl-C`, or the network hangs
- ❖ Sent from kernel (e.g. detects divide by zero (SIGFPE) or could be at the request of another process to send to another)

Signals

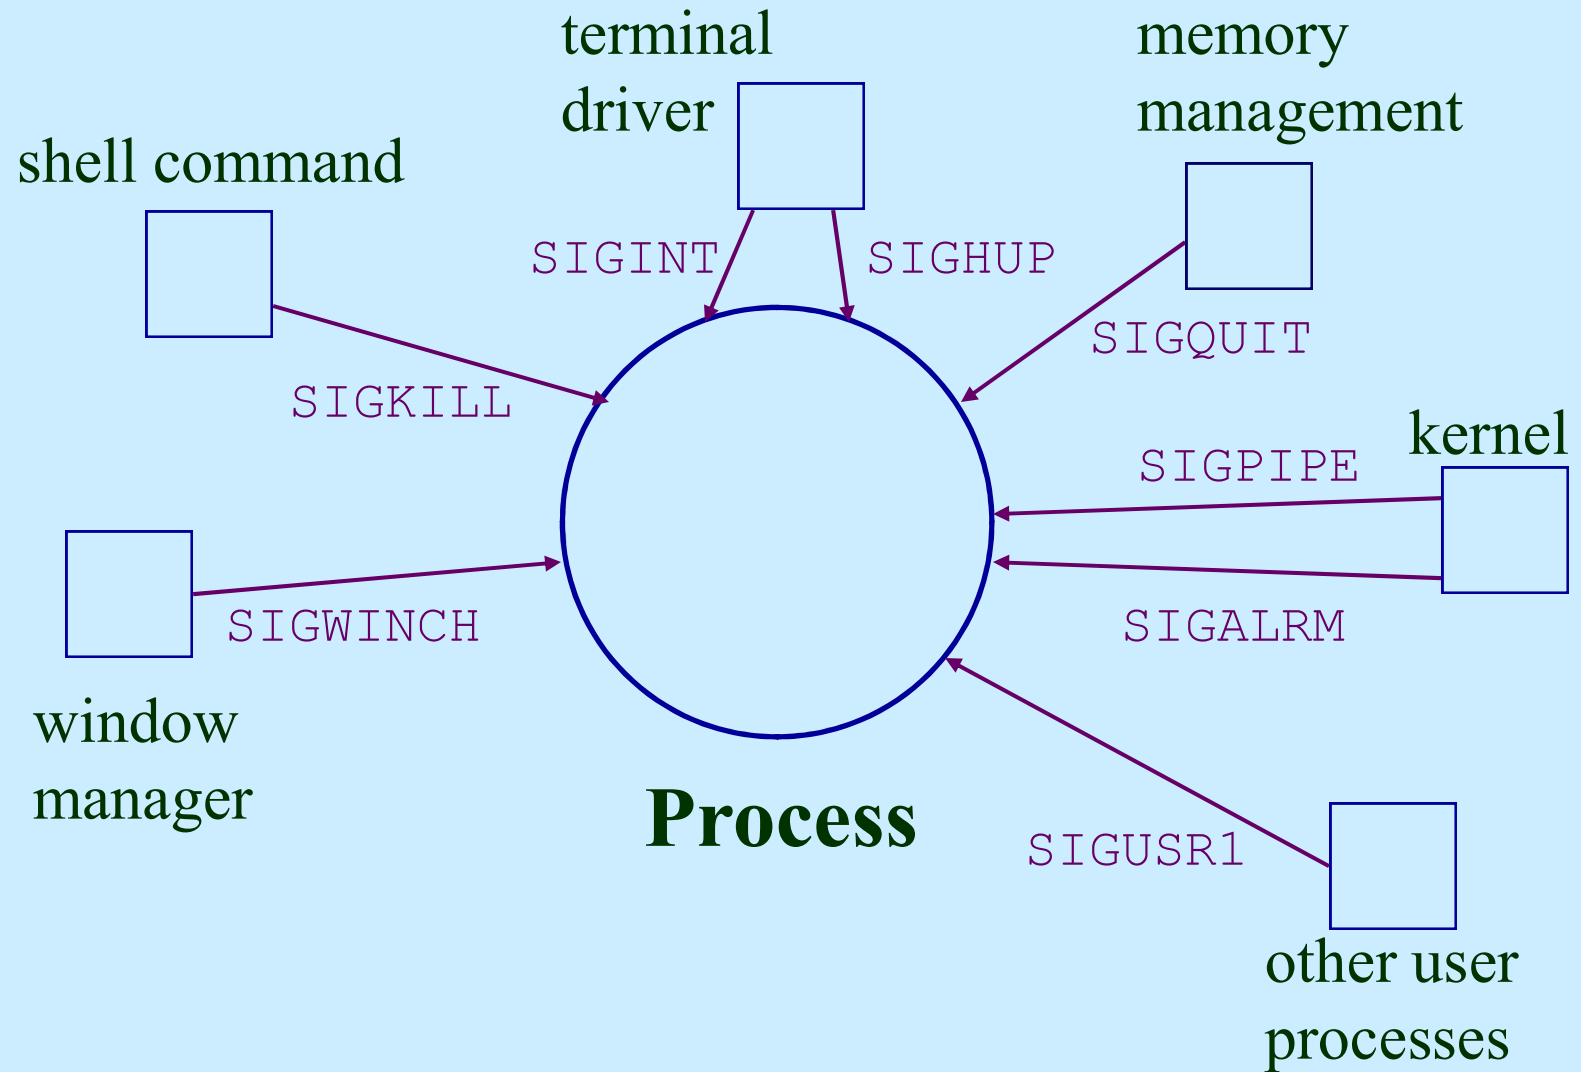
- ❖ Only information that a signal carries is its unique ID and that it arrived
- ❖ Each signal is mapped to a specific bit in the process signal status.
- ❖ Since this is a single bit, only one signal of a specific type can be “seen” by the targeted process. i.e. you can’t stack them !

2. Signal Types (31 in POSIX)

❖ <u>Name</u>	<u>Description</u>	<u>Default Action</u>
SIGINT	Interrupt character typed	terminate process
SIGQUIT	Quit character typed (^\\)	create core image
SIGKILL	kill -9	terminate process
SIGSEGV	Invalid memory reference	create core image
SIGPIPE	Write on pipe but no reader	terminate process
SIGALRM	alarm() clock 'rings'	terminate process
SIGUSR1	user-defined signal type	terminate process
SIGUSR2	user-defined signal type	terminate process

❖ See man 7 signal

Signal Sources



3. Generating a Signal

- ❖ You have probably already sent signals to applications without knowing it.
- ❖ have you pressed CTRL-C (or ^C) to stop a program from running?
- ❖ When you shut down your computer, the O/S sends a shutdown signal to your applications to make sure files are closed properly and there is no data loss (assuming your application can capture the signal).

Generating a Signal

- ❖ Use the UNIX command:

```
$ kill -KILL 4481
```

- send a SIGKILL signal to pid 4481

- check

- ◆ `ps -l`

- to make sure process died

- ❖ `kill` is not a good name; `send_signal` might be better.

kill()

- ❖ Send a signal to a process (or group of processes).

- ❖ `#include <signal.h>`

```
int kill( pid_t pid, int signo );
```

- ❖ Return 0 if ok, -1 on error.

NOTE: the PID must be something you have the “right” to talk to.

Some pid Values

❖ *pid*

Meaning

> 0

send signal to process `pid`

$== 0$

send signal to all processes
whose process group ID
equals the sender's pgid.
e.g. parent kills all children

4. Responding to a Signal

❖ A process can:

- 1) ignore/discard the signal (not possible with SIGKILL or SIGSTOP)
- 2) execute a **signal handler** function, and then possibly resume execution or terminate
- 3) carry out the default action for that signal

❖ The choice is called the process' *signal disposition*

signal(): library call

```
#include <signal.h>
```

```
void (*signal( int signo, void (*func)(int) ))(int);
```

```
typedef void Sigfunc(int); /* my defn */
```

```
Sigfunc *signal( int signo, Sigfunc *handler );
```

- ❖ signal returns a pointer to a function that returns an int (i.e. it returns a pointer to Sigfunc)
- ❖ Specify a signal handler function to deal with a signal type.
- ❖ Returns *previous* signal disposition if ok, SIG_ERR on error.

The signal function itself returns a **pointer** to a function. The return type is the same as the function that is passed in, i.e., a function that takes an int and returns a void

The *handler* function
Receives a single integer as Argument and returns *void*

`#include <signal.h>`

`void (*signal(int sig, void (*handler)(int))) (int) ;`

❖ `signal` returns a pointer to the PREVIOUS signal handler

Signal is a function that takes two arguments:
sig and *handler*

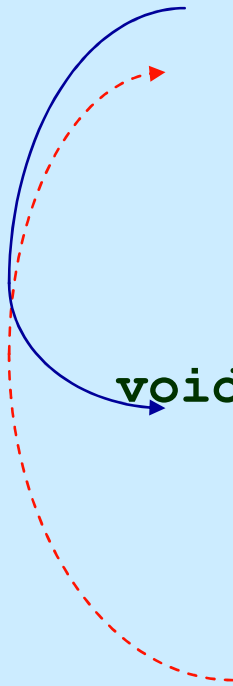
`.h>
func(int) ;
int signo`

The function to be called when the specified signal is received
pointer to *handler*

The returned function takes a integer parameter.

How to Program with Signals

```
int main()  
{  
    signal( SIGQUIT, foo );  
    :  
    /* do usual things until SIGQUIT */  
    return 0;  
}  
  
void foo( int signo )  
{  
    :          /* deal with SIGQUIT signal */  
    return;    /* return to program */  
}
```



sig_examp.c – Signal Handler

```
void sig_usr( int signo )
/* argument is signal number */
{
    if( signo == SIGINT )
        printf("Received ^C SIGINT\n");
    else if( signo == SIGQUIT )
        printf("Received ^\ SIGQUIT\n");
    else {
        printf("Error: received signal %d\n", signo);
        exit(1); }
    return;
}
```

sig_examp.c

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void sig_usr( int signo );    /* handles two signals */

int main()
{
    int i = 0, done = 0;
    if( signal( SIGINT,sig_usr ) == SIG_ERR )
        printf( "Cannot catch ^C SIGINT\n" );

    if( signal( SIGQUIT,sig_usr ) == SIG_ERR )
        printf("Cannot catch ^\ SIGQUIT\n");

    while(!done)
    {
        printf( "%2d\n", i );
        pause();
        /* pause until signal handler
           * has processed signal */
        i++;
    }
    return 0;
}
```


Usage

```
$ ./sig_example
0
^CReceived SIGINT
1
^CReceived SIGINT
2
^\Received SIGQUIT
3
^Z
[1]+  Stopped      ./sig_example
$ kill -9 %1
[1]+  Stopped ./sig_example
```

Special Sigfunc * Values

❖ *Value*

Meaning

SIG_IGN

Ignore / discard the signal.

SIG_DFL

Use default action to handle signal.

SIG_ERR

Returned by `signal()` as an error.

Multiple Signals

- ❖ If many signals of the *same* type are waiting to be handled (e.g. two `SIGQUITs`), then most UNIXs will only deliver *one* of them.
 - the others are thrown away - i.e. pending signals are not queued
 - for each signal type, just have a single bit indicating whether or not the signal has occurred
- ❖ If many signals of *different* types are waiting to be handled (e.g. a `SIGQUIT`, `SIGSEGV`, `SIGUSR1`), they are not delivered in any fixed order.

pause()

- ❖ Suspend the calling process until a signal is caught.
- ❖ `#include <unistd.h>`
`int pause(void);`
- ❖ Returns -1 with `errno` assigned `EINTR`.
- ❖ `pause()` only returns after a signal handler has returned.

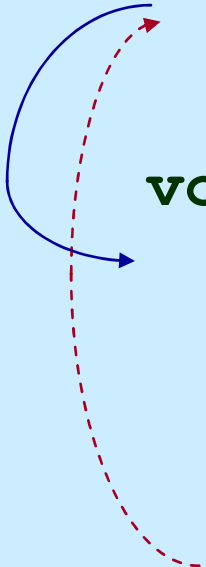
The Reset Problem

- ❖ In many UNIXs, the signal disposition in a process is **reset** to its **default action** immediately after the signal has been delivered.
 - This is a concern with Solaris but not with Linux (loki)
- ❖ Must call `signal()` again to reinstall the signal handler function.

Reset Problem Example

```
int main()
{
    signal(SIGQUIT, foo);
    :
    /* do usual things until SIGQUIT */
}

void foo(int signo)
{
    signal(SIGQUIT, foo); /* reinstall */
    :
    return;
}
```



A diagram consisting of two curved arrows. A solid blue arrow starts from the left side of the `foo` function and points to the `main` function. A dashed red arrow starts from the bottom of the `foo` function and points back to the `main` function, indicating a loop or a return to the caller.

Reset Problem

```
void ouch( int sig )
{
    printf( "OUCH! - I got signal %d\n", sig );
    signal(SIGQUIT, ouch);
}

int main()
{
    int done = 0;
    signal( SIGQUIT, ouch );
    while(!done)
    {
        printf("Hello World!\n");
        sleep(1);
    }
}
```

To keep catching the signal with this function, must call the *signal* system call again (for some versions of Unix)

Problem: from the time that the interrupt function starts to just before the signal handler is re-established the signal will not be handled.

If another SIGQUIT signal is received during this time, default behavior will be done, i.e., program will terminate.

Re-installation may be too slow!

- ❖ There is a (very) small time period in `touch()` when a new `SIGQUIT` signal will cause the default action to be carried out -- process termination.
- ❖ With `signal()` there is no answer to this problem.
 - **POSIX** signal functions **solve** it (and some other later UNIXs)

Another Example

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
int beeps = 0; // GLOBAL VARIABLE !!!
static void handler( int signo )
{
    printf( "BEEP\n" );
    if(++beeps < 5 )
    {
        signal( SIGALRM, handler ); /* may not be needed */
        alarm(1); /* send a SIGALRM in 1 sec */
    }
    else
    {
        printf("BOOM!\n");
        exit(0);
    }
    return;
}
```

```
int main( void )
{
    int i = 0, done = 0;
    if( signal( SIGALRM, handler ) == SIG_ERR )
        printf("Cannot catch SIGALRM\n" );
    alarm(1); /* send a SIGALRM in 1 sec */
    while( !done )
    {
        printf( "%d: ", i );
        pause(); /* wait for any signal */
        i++;
    }
    return 0;
}
```

5. Common Uses of Signals

5.1 Ignore a Signal

5.2 Clean up and Terminate

5.3 Report Status

5.4 Turn Debugging on/off

5.5 Restore Previous Handler

5.1 Ignore a Signal

```
    :  
int main()  
{  
    signal(SIGINT, SIG_IGN);  
    signal(SIGQUIT, SIG_IGN);  
    :  
    /* do work without interruptions */  
}
```

- ❖ Cannot ignore/handle SIGKILL or SIGSTOP
- ❖ Should check for SIG_ERR

5.2 Clean up and Terminate

```

:
/* global variables */
int my_children_pids;
:
void clean_up(int signo);

int main()
{
    signal(SIGQUIT, clean_up);
    :
}

void clean_up(int signo)
{
    unlink("/tmp/work-file");
    kill(my_children_pids, SIGTERM);
    wait((int *)0);
    printf("Program terminated\n");
    exit(1);
}
```

Problems

- ❖ If a program is run in the **background** then the interrupt and quit signals (`SIGINT`, `SIGQUIT`) are automatically ignored.
- ❖ Your code should not override these changes:
 - check if the signal dispositions are `SIG_IGN`

Checking the Disposition

new disposition

old disposition

```
    :  
    if( signal(SIGINT, SIG_IGN) != SIG_IGN )  
        signal(SIGINT, clean_up);  
  
    if( signal(SIGQUIT, SIG_IGN) != SIG_IGN )  
        signal(SIGQUIT, clean_up);  
    :
```

- ❖ *Note:* cannot check the signal disposition without changing it (**sigaction** that we will look at later is different)

5.3 Report Status

```

:
void print_status(int signo);
int count;    /* global */

int main()
{ signal(SIGUSR1, print_status);
  :
  for( count=0; count < BIG_NUM; count++ )
  {
    /* read block from tape */
    /* write block to disk */
  }
  ...
}
```

❖ **Must use global variables for status information**

```
void print_status(int signo)
{
  signal(SIGUSR1, print_status); // depends
  printf("%d blocks copied\n", count);
  return;
}
```


5.4 Turn Debugging on/off

```

:
void toggle_debug(int signo);
int debug = 0; /* initialize here */

int main()
{
    signal(SIGUSR2, toggle_debug);
    /* do work */
    if (debug == 1)
        printf("...");
    ...
}

void toggle_debug(int signo)
{
    signal(SIGUSR2, toggle_debug); // depends
    debug = ((debug == 1) ? 0 : 1);
    return;
}
```

5.5 Restore Previous Handler

```
    :  
    Sigfunc *old_hand;  
  
    /* set action for SIGTERM;  
       save old handler */  
    old_hand = signal(SIGTERM, foobar);  
  
    /* do work */  
  
    /* restore old handler */  
    signal(SIGTERM, old_hand);  
    :
```

6. Implementing a read() Timeout

- ❖ Put an upper limit on an operation that might block forever
 - e.g. `read()`
- ❖ 6.1 `alarm()`
- 6.2 **Bad** `read()` Timeout
- 6.3 `setjmp()` **and** `longjmp()`
- 6.4 **Better** `read()` Timeout

6.1 alarm()

- ❖ Set an alarm timer that will ‘ring’ after a specified number of seconds
 - a SIGALRM signal is generated
- ❖

```
#include <unistd.h>  
long alarm(long secs);
```
- ❖ Returns 0 or number of seconds until previously set alarm would have ‘rung’.

Some Tricky Aspects

- ❖ A process can have **at most one** alarm timer running at once.
- ❖ If `alarm()` is called when there is an **existing** alarm set then it returns the number of seconds remaining for the old alarm, and sets the timer to the new alarm value.
 - What do we do with the “old alarm value”?
- ❖ An `alarm(0)` call causes the previous alarm to be cancelled.

6.2 Bad read() Timeout

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

#define MAXLINE  512

void sig_alm( int signo );

int main()
{
    int n;
    char line[MAXLINE];
    if( signal(SIGALRM, sig_alm) == SIG_ERR )
    {
        printf("signal(SIGALRM) error\n");
        exit(1);
    }
    . . .
```

```
alarm(10);  
n = read( 0, line, MAXLINE );  
alarm(0);  
  
if( n < 0 ) /* read error */  
    fprintf( stderr, "\nread error\n" );  
else  
    write( 1, line, n );  
return 0;  
}
```

```
void sig_alrm(int signo)  
/* do nothing, just handle signal */  
{  
    return;  
}
```

Problems

- ❖ The code assumes that the `read()` call terminates with an error after being interrupted
- ❖ **Race Condition**: The kernel may take longer than 10 seconds to start the `read()` after the `alarm()` call.
 - the alarm may ‘ring’ before the `read()` starts
 - then the `read()` is not being timed; may block forever
 - Two ways to solve this:
 - ◆ **setjmp**
 - ◆ **sigprocmask** and **sigsuspend**

6.3 setjmp() and longjmp()

- ❖ In C we cannot use `goto` to jump to a label in **another** function
 - use `setjmp()` and `longjmp()` for those ‘long jumps’
- ❖ Only uses which are good style:
 - error handling which requires a deeply nested function to recover to a higher level (e.g. back to `main()`)
 - coding timeouts with signals

Nonlocal Jumps: `setjmp()` & `longjmp()`

- ❖ **Powerful (but dangerous) user-level mechanism for transferring control to an arbitrary location**
 - » controlled way to break the procedure call/return discipline
 - » Useful for error recovery and signal recovery
- ❖ **`setjmp(jmp_buf j)`**
 - » called before `longjmp()`
 - » identified return site for subsequent `longjmp()`
 - » On first call, returns a zero
 - » subsequent calls return the value from the int 2nd parameter of `longjump()`
 - » Called once, returns one or more times
- ❖ **Implementation:**
 - » remember where you are by storing the current register context, stack pointer and PC value in `jmp_buf`
 - » returns 0 the first time it is called.

Prototypes

- ❖ `#include <setjmp.h>`
`int setjmp(jmp_buf env);`
`void longjmp(jmp_buf env, int val);`
- ❖ Returns 0 if called directly, non-zero if returning from a call to `longjmp()`.
- ❖ In the `setjmp()` call, `env` is initialized to information about the current state of the stack.
- ❖ The `longjmp()` call causes the stack to be reset to its `env` value.
- ❖ Execution restarts after the `setjmp()` call, but this time `setjmp()` returns `val`.

Example

```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>
void handler(int sig);
jmp_buf buf;

int main()
{
    int done = 0;
    signal(SIGQUIT, handler);
    if( !setjmp(buf) )
        printf("starting\n");
    else
        printf("restarting\n");
    while( !done )
    {
        sleep(1);
        printf("processing...\n");
    }
}
```

```
void handler(int sig)
{
    // may not be needed
    signal(SIGQUIT, handler);
    longjmp( buf, 1 );
}
```

Note: setjmp() returns a zero on first call (FALSE). It returns the value of the 2nd param of longjump() on subsequent invocations.

Output

```
$ ./jumper
starting
processing...
^\\restarting
processing...
^\\restarting
processing...
processing...
^Z
Stopped (user)
$ kill -9 %1
[1] Killed    ./jumper
```

Another Example: sleep1()

```
#include <signal.h>
#include <unistd.h>

void sig_alm( int signo )
{
    return; /* return to wake up pause */
}

unsigned int sleep1( unsigned int nsecs )
{
    if( signal( SIGALRM, sig_alm ) == SIG_ERR )
        return (nsecs);
    alarm( nsecs );           /* starts timer */
    pause();                  /* next caught signal wakes */
    return( alarm( 0 ) );     /* turn off timer, return unslept
                               * time */
}
```

Comments on sleep(1)

❖ Alarm erases “old” set alarms set by process

- » Look at return value from the previous alarm() call
 - ◆ If less than new alarm() - wait until old alarm() expires
 - ◆ If larger than new alarm()- reset old alarm() with remaining seconds when done with new alarm()

❖ Race condition

- » between first call to alarm and the call to pause
- => never get out of pause (fix via setjmp/longjmp or sigprocmask/sigsuspend)

sleep2(): Fixes race condition

```
static void jmp_buf env_alm;  
void sig_alm( int signo )  
{  
    longjmp( env_alm, 1 );  
}  
unsigned int sleep2( unsigned int nsecs )  
{  
    if( signal( SIGALRM, sig_alm ) == SIG_ERR )  
        return (nsecs);  
    if( !setjmp( env_alm) ) /* returns 0 (FALSE) on 1st call */  
    {  
        alarm( nsecs );      /* starts timer */  
        pause();             /* next caught signal wakes */  
    }  
    return( alarm( 0 ) );  
}
```


Sleep1 and Sleep2

- ❖ Sleep2 fixes race condition. Even if the pause is never executed.
 - » **A SIGALRM causes sleep2() to return**
 - » **Avoids entering pause() via longjmp()**
- ❖ There is one more problem
 - » **SIGALRM could interrupt some other signal handler and subsequently abort it by executing the longjmp**

6.4 Better read() Timeout

```
#include <stdio.h>
#include <unistd.h>
#include <setjmp.h>
#include <signal.h>

#define MAXLINE  512
void sig_alm( int signo );
jmp_buf env_alm;

int main()
{ int n;
  char line[MAXLINE];

  if( signal(SIGALRM, sig_alm) == SIG_ERR)
  {
    printf("signal(SIGALRM) error\n");
    exit(1);
  }
```

```
if( setjmp(env_alm) != 0 )
{
    printf("\nread() too slow\n");
    exit(2);
}
```

```
alarm(10);
n = read(0, line, MAXLINE);
alarm(0);

if( n < 0 ) /* read error */
    printf("\nread error\n" );
else
    write( 1, line, n );
return 0;
}
```

```
void sig_alm(int signo)
/* interrupt the read() and jump to setjmp() call with value 1 */
{
    longjmp(env_alm, 1);
}
```

Caveat: Non-local Jumps

From the UNIX **man** pages:

WARNINGS

If `longjmp()` or `siglongjmp()` are called even though `env` was never primed by a call to `setjmp()` or `sigsetjmp()`, or when the last such call was in a function that has since returned, absolute chaos is guaranteed.

A Problem Remains!

- ❖ If the program has several signal handlers then:
 - execution might be inside one when an alarm ‘rings’
 - the `longjmp()` call will jump to the `setjmp()` location, and abort the other signal handler -- might lose / corrupt data

Stop here Winter 2023

7. POSIX Signal Functions

- ❖ The POSIX signal functions can control signals in more ways:
 - can *block signals* for a while, and deliver them later (good for coding critical sections)
 - can *switch off the resetting* of the signal disposition when a handler is called (no reset problem)
- ❖ The POSIX signal system, uses **signal sets**, to deal with pending signals that might otherwise be missed while a signal is being processed

7.1 Signal Sets

- ❖ The signal set stores collections of signal types.
- ❖ Sets are used by signal functions to define which signal types are to be processed.
- ❖ POSIX contains several functions for creating, changing and examining signal sets.

Prototypes

❖ `#include <signal.h>`

```
int sigemptyset( sigset_t *set );
```

```
int sigfillset( sigset_t *set );
```

```
int sigaddset( sigset_t *set, int signo );
```

```
int sigdelset( sigset_t *set, int signo );
```

```
int sigismember( const sigset_t *set, int signo );
```

Cont'd

- ❖ *sigemptyset* - initializes signal set pointed by set so that all signals are excluded
- ❖ *sigfillset* - all signals are included
 - One of *sigemptyset* or *sigfillset* must be called to initialize a signal set
- ❖ *sigaddset* - add a single signal (signo) to set
- ❖ *sigdelset* - remove signo from set
- ❖ *sigismember* - is signo from a member of set

7.2 sigprocmask()

- ❖ A process uses a signal set to create a mask which defines the signals it is **blocking** from delivery.
 - good for critical sections where you want to block certain signals.

```
#include <signal.h>
int sigprocmask( int how, const sigset_t *set,
                 sigset_t *oldset);
```

- ❖ *how* – indicates how mask is modified

how Meanings

❖ *Value* *Meaning*

SIG_BLOCK set signals are **added** to mask

SIG_UNBLOCK set signals are **removed** from mask

SIG_SETMASK set **becomes new** mask

A Critical Code Region

```
sigset_t newmask, oldmask;

sigemptyset( &newmask );
sigaddset( &newmask, SIGQUIT );

/* block SIGQUIT; save old mask */
sigprocmask( SIG_BLOCK, &newmask, &oldmask );

/* critical region of code */

/* reset mask which unblocks SIGQUIT */
sigprocmask( SIG_SETMASK, &oldmask, NULL );
```

7.3 sigaction()

- ❖ `#include <signal.h>`

```
int sigaction(int signo, const struct sigaction *act,  
              struct sigaction *oldact );
```

- ❖ Supercedes (more powerful than) `signal()`
 - `sigaction()` can be used to code a non-resetting `signal()`
- ❖ ***signo*** is the signal you want to perform an action on
- ❖ ***act*** is the action
- ❖ ***oact*** is the old action (can be set to `NULL`, if uninteresting)
- ❖ Cannot handle `SIGSTOP` and `SIGKILL`

sigaction() Structure

```
struct sigaction
{
    void      (*sa_handler)( int );
               /* action to be taken or SIG_IGN, SIG_DFL */
    sigset_t  sa_mask; /* additional signals to be blocked */
    int       sa_flags; /* modifies action of the signal */
}
```

❖ sa_flags –

- SA_RESETHAND: reset handler to default upon return
- SA_RESTART: if process is interruptible, then it is restarted

sigaction() Behavior

- ❖ A *signo* signal causes the *sa_handler* signal handler to be called.
- ❖ **While *sa_handler* executes**, the signals in *sa_mask* are blocked. Any additional *signo* signals are also blocked.
- ❖ *sa_handler* remains installed until it is changed by another `sigaction()` call. **No reset problem.**

Signal Raising

```
#include <signal.h>
#include <stdio.h>
void ouch( int signo )
{ printf( "OUCH! signo = %d\n", signo ); }
```

```
struct sigaction
{
    void (*) (int) sa_handler
    sigset_t sa_mask
    int sa_flags
}
```

```
int main()
{
    int done = 0;
    struct sigaction act;
    act.sa_handler = ouch;
    sigemptyset( &act.sa_mask );
    act.sa_flags = 0;
    sigaction( SIGQUIT, &act, 0 );
    printf( "Hello World!\n" );
}
```

Set the signal handler to be the function ouch

No flags are needed here. Possible flags include: SA_RESETHAND

We can manipulate sets of signals to be blocked while *signo* is processed.

This call sets the signal handler for the SIGQUIT (Ctrl-\) signal

Signal Raising

- ❖ This function will continually capture the Ctrl-\ (SIGQUIT) signal
- ❖ Default behavior is **not** restored after signal is caught.
- ❖ To terminate the program, must type Ctrl-c (SIGINT), or Ctrl-z (SIGSTOP) followed by kill %1

sigexPOS.c

```
/* include files as before */
#include <signal.h>
int main(void)
{
    /* struct to deal with action on signal set */
    static struct sigaction act;
    void catchquit(int); /* user signal handler */
    /* set up action to take on receipt of SIGQUIT */
    act.sa_handler = catchquit;
    /* create full set of signals */
    sigfillset(&act.sa_mask);
    /* before sigaction, SIGQUIT will terminate */
    /* now, SIGQUIT will cause catchquit to execute */
    sigaction( SIGQUIT, &act, NULL );
}
```

```

/* remainder of main program */
printf("sleep call #1\n");
sleep(3);
printf("sleep call #2\n");
sleep(3);
printf("sleep call #3\n");
sleep(3);
printf("sleep call #4\n");
sleep(3);
printf("Exiting \n");
exit (0);
}

/* simple signal handler */
void catchquit(int signo)
{
    printf("\nCATCHQUIT: signo=%d\n", signo);
    printf("\nCATCHQUIT: returning\n\n");
}

```

```

$ ./sleeper
Sleep call #1
Sleep call #2
^
CATCHQUIT: signo = 3
CATCHQUIT: returning
Sleep call #3
Sleep call #4
Exiting
$

```

Signals - Ignoring signals

- ❖ Other than SIGKILL and SIGSTOP, signals can be ignored:
- ❖ Instead of in the previous program:

```
act.sa_handler = catchquit /* or whatever */
```

We use:

```
act.sa_handler = SIG_IGN;
```

The ^\ key will be ignored

Restoring previous action

- ❖ The third parameter to `sigaction`, `oact`, can be used:

```
/* save old action */  
sigaction( SIGTERM, NULL, &oact );
```

```
/* set new action */  
act.sa_handler = SIG_IGN;
```

```
sigaction( SIGTERM, &act, NULL );
```

```
/* restore old action */  
sigaction( SIGTERM, &oact, NULL );
```

7.4 Other POSIX Functions

- ❖ `sigpending()` examine blocked signals
- ❖ `sigsetjmp()`
 `siglongjmp()` jump functions for use
 in signal handlers which
 handle masks correctly
- ❖ `sigsuspend()` atomically reset mask
 and sleep

[sig]longjmp & [sig]setjmp

NOTES (longjmp, sigjmp)

POSIX does not specify whether longjmp will restore the signal context. If you want to save and restore **signal masks**, use siglongjmp.

NOTES (setjmp, sigjmp)

POSIX does not specify whether setjmp will save the signal context. (In SYSV it will not. In BSD4.3 it will, and there is a function _setjmp that will not.) If you want to save signal masks, use sigsetjmp.

Example

```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>

sigjmp_buf buf;

void handler(int sig)
{
    siglongjmp(buf, 1);
}

main()
{
    int done = 0;
    signal(SIGQUIT, handler);

    if( sigsetjmp(buf, 1) == 0 )
        printf("starting\n");
    else
        printf("restarting\n");
    ...
}
```

```
...
while(!done)
{
    sleep(5);
    printf("  waiting...\n");
}
}
```

```
> ./longjump
starting
  waiting...
  waiting... ← Control-\
restarting
  waiting...
  waiting...
  waiting...
restarting ← Control-\
  waiting...
restarting ← Control-\
  waiting...
  waiting...
```

8. Interrupted System Calls

- ❖ When a system call (e.g. `read()`) is interrupted by a signal, a signal handler is called, returns, and then what?
- ❖ On many UNIXs, *slow* system function calls do not resume. Instead they return an error and `errno` is assigned `EINTR`.

Slow System Functions

- ❖ Slow system functions carry out I/O on things that can possibly block the caller forever:
 - pipes, terminal drivers, networks
 - some IPC functions
 - `pause()`, some uses of `ioctl()`

Non-slow System Functions

- ❖ Most system functions are non-slow, including ones that do *disk* I/O
 - e.g. `read()` of a disk file
 - `read()` is sometimes a slow function, sometimes not
- ❖ Some UNIXs restart non-slow system functions after the handler has finished.
- ❖ Some UNIXs only call the handler after the non-slow system function call has finished.

9. System Calls inside Handlers

- ❖ If a system function is called inside a signal handler then it may interact with an interrupted call to the same function in the main code.
 - e.g. `malloc()`
- ❖ This is not a problem if the function is *reentrant*
 - a process can contain multiple calls to these functions at the same time
 - e.g. `read()`, `write()`, `fork()`, many more

Non-reentrant Functions

- ❖ A functions may be non-reentrant (only one call to it at once) for a number of reasons:
 - it uses a static data structure
 - it manipulates the heap: `malloc()`, `free()`, etc.
 - it uses the standard I/O library
 - ◆ e.g, `scanf()`, `printf()`
 - ◆ the library uses global data structures in a non-reentrant way