

UNIX Systems Programming

Interprocess communication



Overview

1. What is a Pipe
2. Unix System Review
3. Processes (review)
4. Pipes
5. FIFOs
6. Shared Memory
7. Sockets

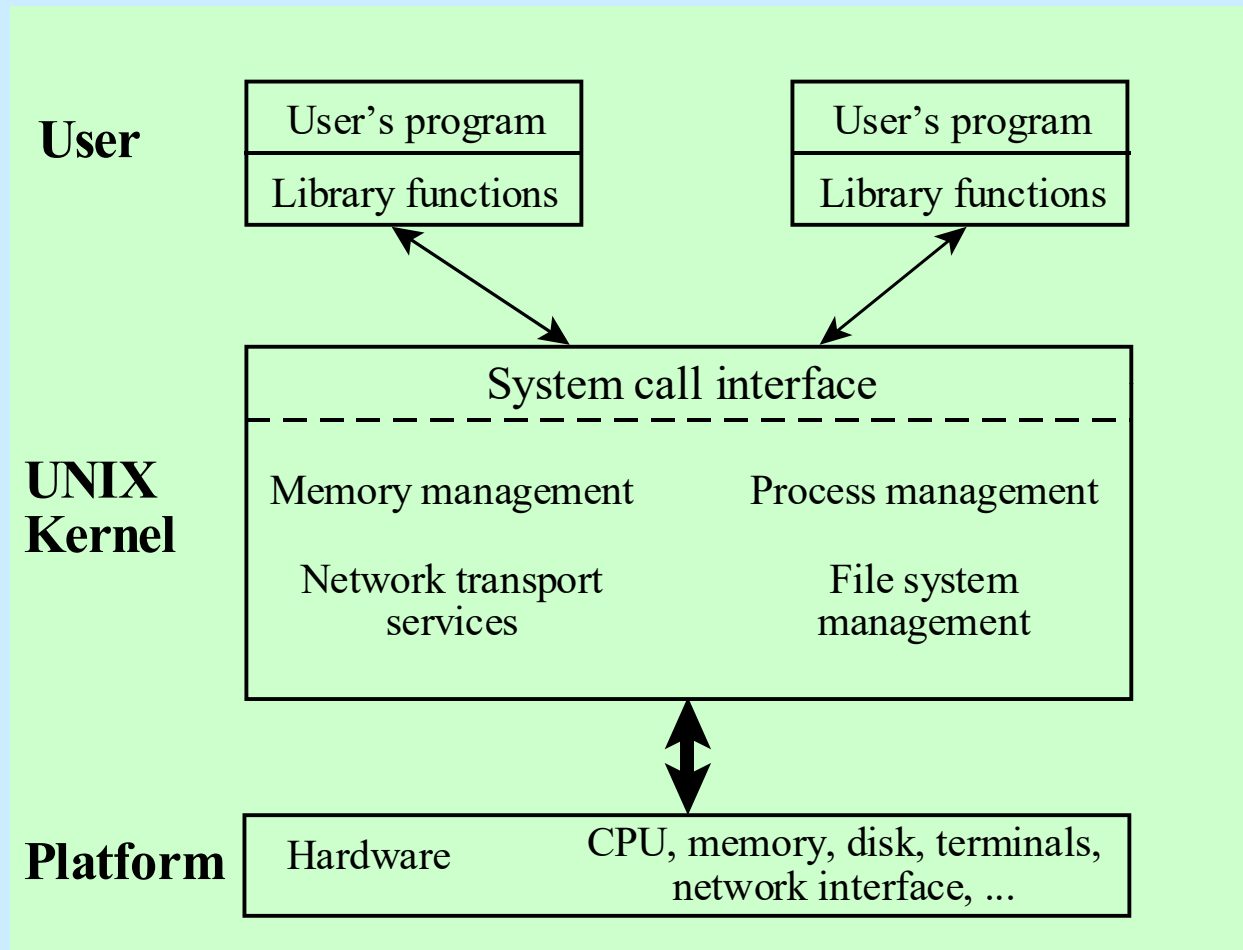
1. Pipes

- A form of interprocess communication between processes that have a common ancestor
- It is a **one-way** (half duplex) communication channel which can be used to link processes
- A pipe is a generalization of the file idea
 - Can use I/O functions like `read()` and `write()` to receive and send data
- Typical use:
 - Pipe created by a process
 - Process calls `fork()`
 - Pipe used between parent and child

Differences between versions

- All systems support half-duplex
 - Data flows in only one direction
- Many newer systems support full duplex
 - Data flows in two directions
- For portability, assume only half-duplex
- Pipes at the UNIX shell level
 - `who | wc -1`
 - gives a count of the number of users logged on

A UNIX System



Review: fork()

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork( void );
```

- Creates a child process by making a *copy* of the parent process
- Both the child *and* the parent continue running

Context used by child and exec()

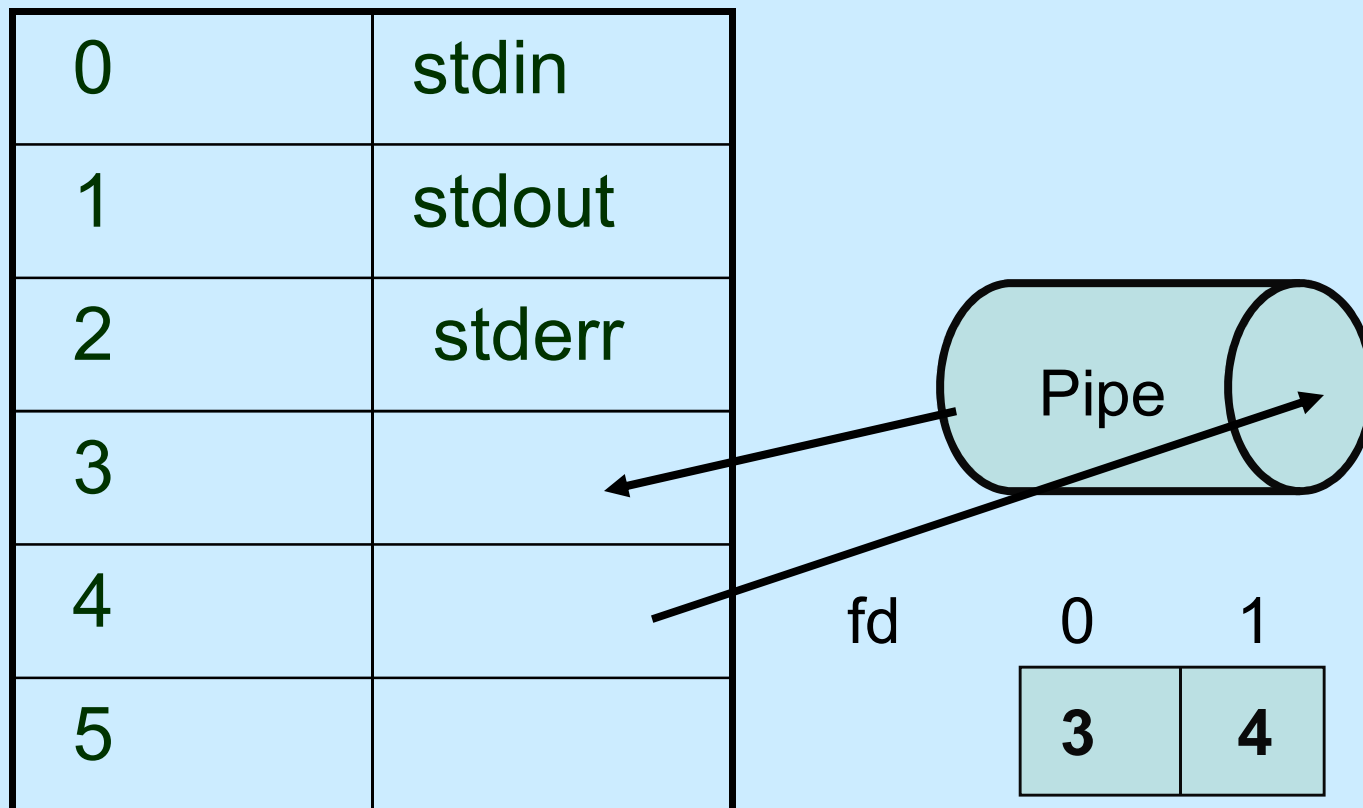
Attribute	Inherited by child	Retained in exec()
PID	No	Yes
Real PID	Yes	Yes
Effective PID	Yes	Depends on <i>setuid</i> bit
Data	Copied	No
Stack	Copied	No
Heap	Copied	No
Program Code	Shared	No
File Descriptors	Copied (file ptr is shared)	Usually
Environment	Yes	Depends on <i>exec()</i>
Current Directory	Yes	Yes
Signal	Copied	Partially

Programming with Pipes

```
#include <unistd.h>
int pipe(int fd[2]);
```

- Returns 0 if ok, -1 on error
- Pipe() binds fd[] with two file descriptors
 - fd[0] is open for reading
 - fd[1] is open for writing
 - Output of fd[1] is input to fd[0]


After the pipe() call



Example: pipe-yourself.c

```
#include <stdio.h>
#include <unistd.h>
#define MSGSIZE 16 /* null */
char *msg1="hello, world #1";
char *msg2="hello, world #2";
char *msg3="hello, world #3";
int main()
{
    char inbuf[MSGSIZE];
    int p[2], i;
    if( pipe( p ) < 0 )
        { /* open pipe */
            perror( "pipe" );
            exit( 1 ); }
    write( p[1], msg1, MSGSIZE );
    write( p[1], msg2, MSGSIZE );
    write( p[1], msg3, MSGSIZE );
    for( i=0; i < 3; i++ )
        { /* read pipe */
            read( p[0], inbuf, MSGSIZE );
            printf( "%s\n", inbuf ); }
    return 0;
}
```

```
$ a.out
hello, world #1
hello, world #2
hello, world #3
$
```



Things to Note

- Pipes uses FIFO ordering: *first-in first-out*.
- Read/write amounts **do not** need to be the same, but then text will be split differently.
- Pipes are most useful with `fork()` which creates an IPC connection between the parent and the child (or between the parents children)


Example: pipe_fork.c

```
#include <stdio.h>
#include <sys/wait.h>
#include <unistd.h>
#define MSGSIZE 16
char *msg1="hello, world #1";
char *msg2="hello, world #2";
char *msg3="hello, world #3";
int main()
{
    char inbuf[MSGSIZE];
    int p[2], i, pid;
    if( pipe( p ) < 0 )
        { /* open pipe */
            perror( "pipe" );
            exit( 1 );
        }
    if( (pid = fork()) < 0 )
        {
            perror( "fork" );
            exit( 2 );
        }
}
```

Cont'd

```
else if( pid > 0 ) /* parent */
{
    write( p[1], msg1, MSGSIZE );
    write( p[1], msg2, MSGSIZE );
    write( p[1], msg3, MSGSIZE );
    wait( (int *) 0 );
}
else if( pid == 0 ) /* child */
{
    for( i=0; i < 3; i++ )
    {
        read( p[0], inbuf, MSGSIZE );
        printf( "%s\n", inbuf );
    }
}
return 0;
}
```


```
$ a.out
hello, world #1
hello, world #2
hello, world #3
$
```



Another look


0	stdin
1	stdout
2	stderr
3	
4	
5	

Parent



0	stdin
1	stdout
2	stderr
3	
4	
5	

Child



Things to Note

- Notice that both parent and child can read/write to the pipe
- Possible to have multiple readers/writers attached to a pipe
 - Can causes confusion
- Best style is to **close** links you do not need
 - i.e, we close the read end in one process and the write end in the other process
 - For our example, the read end of the parent and the write end of the child


Example: pipe_fork_close.c

```
#include <stdio.h>
#include <sys/wait.h>
#include <unistd.h>
#define MSGSIZE 16
char *msg1="hello, world #1";
char *msg2="hello, world #2";
char *msg3="hello, world #3";
int main()
{
    char inbuf[MSGSIZE];
    int p[2], i, pid;
    if( pipe( p ) < 0 )
        { /* open pipe */
            perror( "pipe" );
            exit( 1 );
        }
    if( (pid = fork()) < 0 )
        {
            perror( "fork" );
            exit( 2 );
        }
```


Cont'd

```
else if( pid > 0 ) /* parent */
{
    close( p[0] ); /* read link */
    write( p[1], msg1, MSGSIZE );
    write( p[1], msg2, MSGSIZE );
    write( p[1], msg3, MSGSIZE );
    wait( (int *) 0 );
}
else if( pid == 0 ) /* child */
{
    close( p[1] ); /* write link */
    for( i=0; i < 3; i++ )
    {
        read( p[0], inbuf, MSGSIZE );
        printf( "%s\n", inbuf );
    }
}
return 0;
}
```


```
$ a.out
hello, world #1
hello, world #2
hello, world #3
$
```



Another look

0	stdin
1	stdout
2	stderr
3	X
4	
5	

Parent



0	stdin
1	stdout
2	stderr
3	
4	X
5	

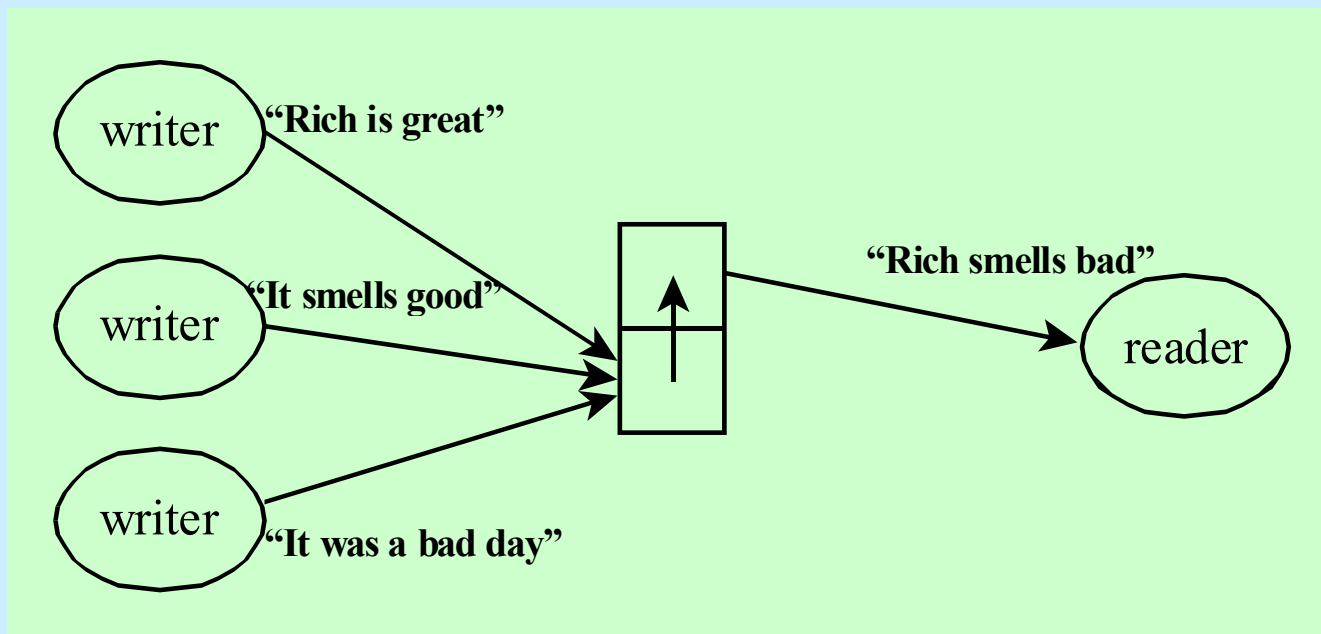
Child

Rules of Pipes

- Every pipe has a size limit
 - POSIX minimum is 512 bytes (most systems makes this figure larger ... for Solaris it is 5120 bytes)
- `read()` blocks if pipe is empty *and* there is a write link open to that pipe
 - Close write links or `read()` will never return
- `read()` from a pipe whose `write()` end is closed *and* is empty returns 0 (indicates EOF)
- `write()` to a pipe with no `read()` ends returns -1 and generates SIGPIPE and `errno` is set to EPIPE
- `write()` blocks if the pipe is full or there is not enough room to support the `write()` .
 - May block in the middle of a `write()`

Several Writers

- Since a `write()` can suspend in the middle of its output then output from multiple writers may be **mixed up** (*interleaved*).



- In `limits.h`, the constant `PIPE_BUF` (512–4096) gives the maximum number of bytes that can be output by a `write()` without any chance of interleaving
- Use `PIPE_BUF` if there are to be multiple writers in your code

Non-blocking read() & write()

- **Problem:**
 - Sometimes you want to prevent `read()` and `write()` from blocking.
- **Goals:**
 - want to return an error code instead
 - want to poll several pipes in turn until one has data
- **Approaches:**
 - Use `fstat()` on the pipe to get the number of characters in pipe (caveat: multiple readers may give a race condition)
 - Use `fcntl()` on the pipe and set it to `O_NONBLOCK`


Using `fcntl()`

```
#include <sys/types.h>
#include <unistd.h>
#include <fcntl.h>
:
if( fcntl( fd, F_SETFL, O_NONBLOCK ) < 0 )
    perror("fcntl");
:
```

- **Non-blocking write:** On a write-only file descriptor, `fd`, future writes will never block
 - Instead return immediately with a -1 and set `errno` to `EAGAIN`
- **Non-blocking read:** On a read-only file descriptor, `fd`, future reads will never block
 - return -1 and set `errno` to `EAGAIN` or return 0 if pipe is empty (or closed)

Example: Non-blocking with -1 return

- Child writes “hello” to parent every 3 seconds (3 times).
- Parent does a non-blocking read each second.



Example: pipe_nonblocking.c

```
#include <unistd.h>
#include <fcntl.h>
#include <errno.h>
#include <stdio.h>
#include <stdlib.h>
#define MSGSIZE 6
char *msg1="hello";

void parent_read( int p[] );
void child_write( int p[] );

int main()
{
    int pfd[2];
    if( pipe( pfd ) < 0 )
        { /* open pipe */
            perror( "pipe" );
            exit( 1 );
        }
}
```


main Cont'd

```
if( fcntl( pfd[0], F_SETFL, O_NONBLOCK ) < 0 )
    { /* read non-blocking */
        perror( "fcntl" );
        exit( 2 );
    }
switch( fork() )
{
    case -1: /* error */
        perror("fork" );
        exit(3);
    case 0: /* child */
        child_write( pfd );
        break;
    default: /* parent */
        parent_read( pfd );
        break;
}
return 0;
}
```

void parent_read()

```
void parent_read( int p[] )
{
    int nread, done = 0;
    char buf[MSGSIZE];
    close( p[1] ); /* write link */
    while( !done )
    {
        nread = read( p[0], buf, MSGSIZE );
        switch( nread )
        {
            case -1:
                if( errno == EAGAIN )
                {
                    printf( "(pipe empty) \n" );
                    sleep( 1 );
                    break;
                }
            }
        }
    }
}
```

Cont'd

```
        else
        {
            perror( "read" );
            exit(4);
        }

    case 0:
        /* pipe has been closed */
        printf( "End conversation\n" );
        close( p[0] ); /* read fd */
        exit(0);
    default: /* text read */
        printf( "MSG=%s\n", buf );
    } /* switch */
} /* while */
} /* parent_read */
```

void child_write()

```
void child_write( int p[] )
{
    int i;
    close( p[0] ); /* read link */
    for( i = 0; i < 3; i++ )
    {
        write( p[1], msg1, MSGSIZE );
        sleep( 3 );
    }
    close( p[1] ); /* write link */
}
```

```
$ a.out
MSG=hello
(pipe empty)
(pipe empty)
(pipe empty)
MSG=hello
(pipe empty)
(pipe empty)
(pipe empty)
MSG=hello
(pipe empty)
(pipe empty)
(pipe empty)
End conversation
$
```

Limitations of Pipes

- **Processes using a pipe must come from a common ancestor:**
 - e.g. parent and child
 - cannot create general servers like print spoolers or network control servers since unrelated processes cannot use it
- **Pipes are not permanent**
 - they disappear when the process terminates
- **Pipes are sometimes one-way:**
 - makes fancy communication harder to code
- **Pipes do not work over a network**

What are FIFOs/Named Pipes?

- Similar to pipes (as far as `read/write` are concerned, e.g. FIFO channels), but with some additional advantages:
 - Unrelated processes can use a FIFO.
 - A FIFO can be created separately from the processes that will use it.
 - FIFOs look like files:
 - have an owner, size, access permissions
 - open, close, delete like any other file
 - **permanent** until deleted with `rm`

Creating a FIFO

- **UNIX `mkfifo` command:**

```
$ mkfifo fifo1
```

- **On older UNIXs (original ATT UNIX), use `mknod`:**

```
$ mknod fifo1 p
```

- **Use `ls` to get information:**

```
$ ls -l fifo1
```

```
prw----- 1 rhurley staff 0 Jul 3 12:02 fifo1
```

Using FIFOs: FIFO Blocking

- FIFOs can be read and written using standard UNIX commands connected via “<” and “>” (a commands input or output)
- If there are no writers then a **read**:
 - e.g. `cat < fifo1`
will block until there is 1 or more writers.
- If there are no readers then a **write**:
 - e.g. `ls -l > fifo1`
will block until there is 1 or more readers

Reader / Writer Example

```
$ cat < fifo1 &  
[1] 22341  
$ ls -l > fifo1; wait  
total 17  
prw----- 1 rhurley staff 0 Jul 3 12:15 fifo1  
[1] Done cat < fifo1  
$
```

1. Output of `ls -l` is written down the FIFO
2. Waiting `cat` reads from the FIFO and display the output
3. `cat` exits since `read` returns 0 (the FIFO is not open for writing anymore and 0 is returned as EOF)

wait - causes the shell to wait until `cat` exits before redisplaying the prompt

Creating a FIFO in C

```
#include <sys/types.h>
#include <sys/stat.h>
int mkfifo(const char *pathname, mode_t mode);
```

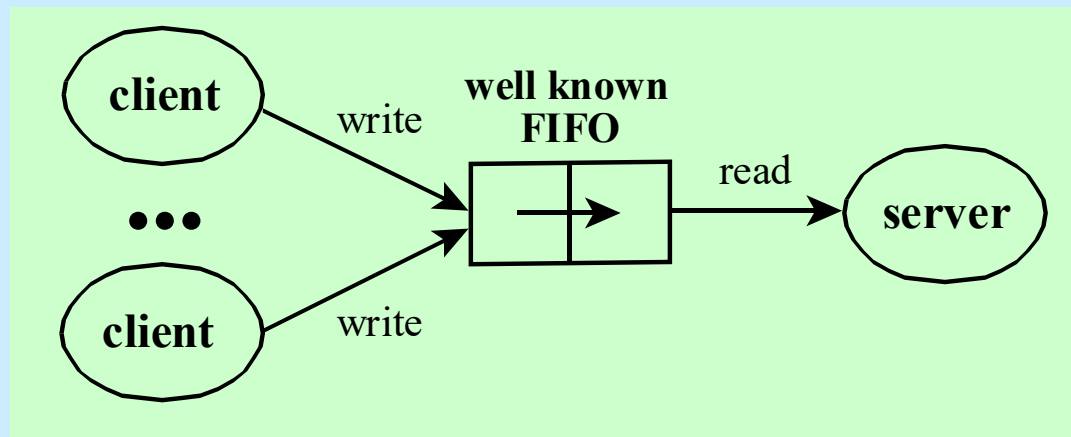
- **Returns 0 if OK, -1 on error.**
- **mode is the same as for open () - and is modifiable by the process' umask value**
- **Once created, a FIFO must be opened using open ()**

Two Main Uses of FIFOs

1. **Used by shell commands to pass data from one shell pipeline to another without using temporary files.**
2. **Create client-server applications on a single machine.**

Client-Server Applications

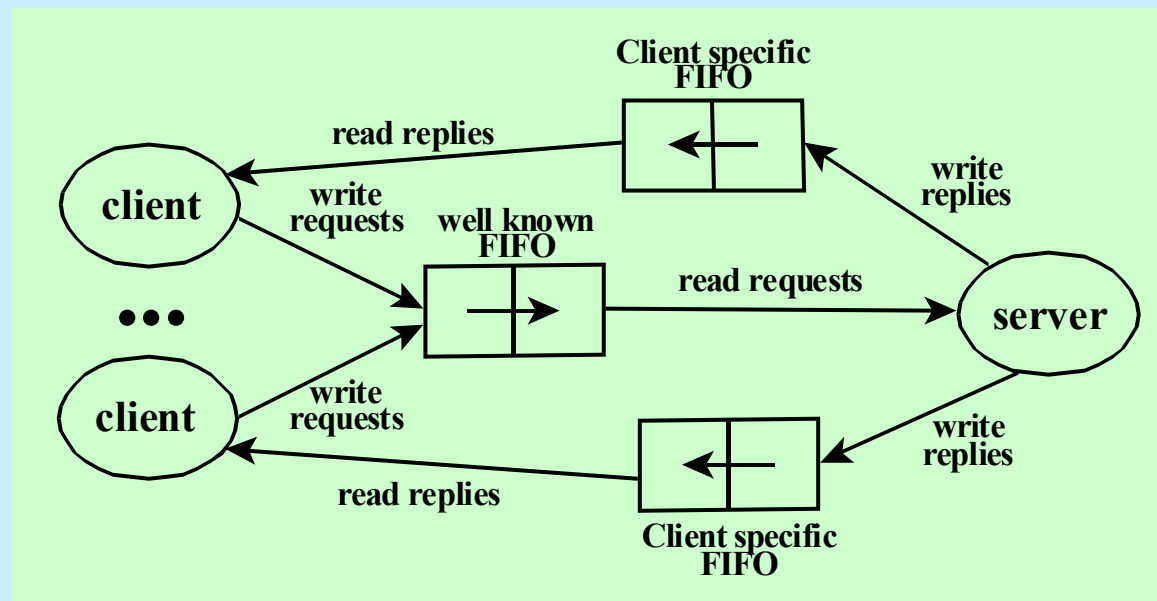
- **Server contacted by numerous clients via a well-known FIFO**



- **How are replies from the server sent back to each client?**

Client-Server FIFO Application

- **Problem:** A single FIFO (as before) is not enough.
- **Solution:** Each client send its PID as part of its message. Which it then uses to create a special 'reply' FIFO for each client
 - e.g. /tmp/serv1.xxxx where xxxx is replaced with the clients process ID



Problems

- **The server does not know if a client is still alive**
 - may create FIFOs which are never used
 - client terminates before reading the response (leaving FIFO with one writer and no reader)
- **Each time number of clients goes from 1 client to 0 the server reads an EOF on the well-known FIFO, if it is set to read-only.**
 - Common trick is to have the server open the FIFO as read-write

Programming Client-server Applications

- First we must see how to create, open and read a FIFO from within C.
- Clients will write in **non-blocking** mode, so they do not have to wait for the server process to start.

Creating a FIFO

```
#include <sys/types.h>
#include <sys/stat.h>
:
int mkfifo(const char *pathname, mode_t mode);
```

- **Creates a FIFO file named by pathname**
- **The FIFO will be given mode permissions (0666)**
- **Can be modified using the process' umask value**

Opening FIFOs

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
:
fd = open( "fifo1", O_WRONLY );
:
```

- **A FIFO can be opened with `open()` (most I/O functions work with pipes).**

Blocking open ()

- An open () call for *writing* will block until another process opens the FIFO for *reading*.
 - this behavior is not suitable for a client who does not want to wait for a server process before sending data.
- An open () call for *reading* will block until another process opens the FIFO for *writing*.
 - this behavior is not suitable for a server which wants to poll the FIFO and continue if there are no readers at the moment.

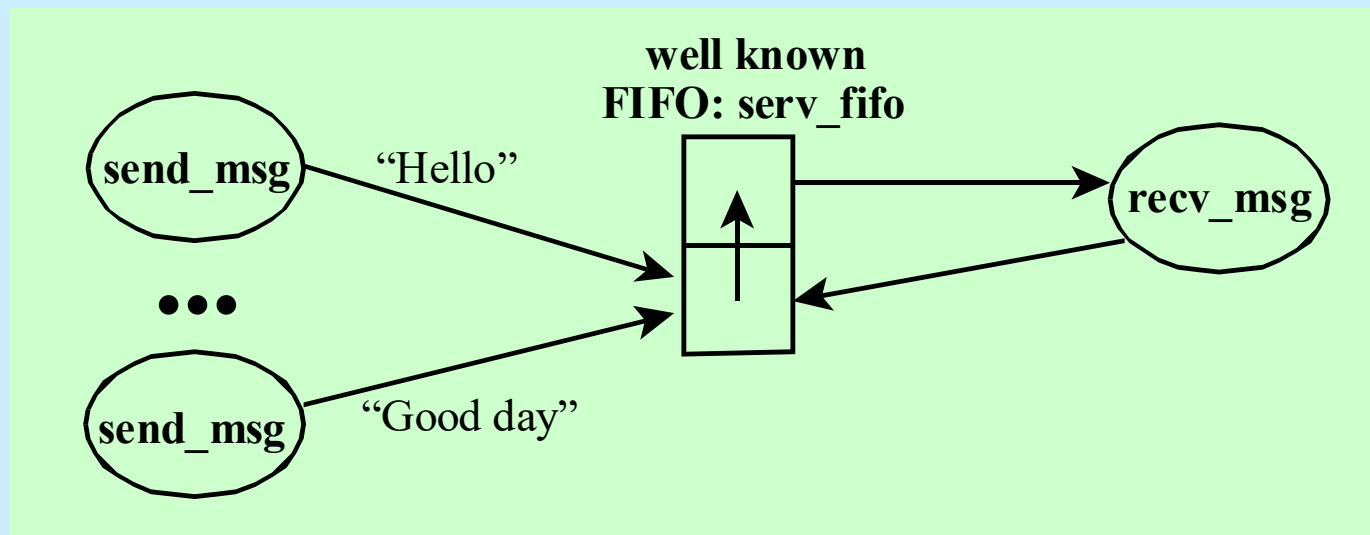
Non-blocking open ()

```
if ( fd = open( "fifo1", O_WRONLY | O_NONBLOCK) ) < 0 )  
    perror( "open FIFO" );
```

- opens the FIFO for writing
- returns -1 and errno is set to ENXIO if there are **no readers**, instead of blocking.
- Later write () calls will also not block.

Example: send_msg, recv_msg

- implement a message system
- exploits the fact that reads/writes to pipes/FIFOs are **atomic**
- if fixed-sized messages are passed, individual messages will stay intact even with **concurrent** senders



Notes:

- `recv_msg` can **read and write**;
 - otherwise the program would block at the open call
 - also avoids responding to reading a “return of 0” when the number of `send_msg` processes goes from 1 to 0 (and the FIFO is empty) `O_RDWR` - ensures that at least one process has the FIFO open for writing (i.e. `recv_msg` itself) so read will always block until data is written to the FIFO
- `send_msg` sends **fixed-size** messages of length `PIPE_BUF` to avoid interleaving problems with other `send_msg` calls. It uses non-blocking.
- `serv_fifo` is globally known, and previously created with `mkfifo`

Header for files

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <fcntl.h>
#include <string.h>
#include <limits.h>
#define SF "serv_fifo"
#define PIPE_BUF 1024
```

send_msg.c

```
void make_msg( char mb[], char input[] );
int main( int argc, char *argv[] )
{
    int fd, i;
    char msgbuf[PIPE_BUF];
    if( argc < 2 )
    {
        printf( "Usage: send-msg msg...\n" );
        exit( 1 );
    }
    if( (fd = open( SF, O_WRONLY | O_NONBLOCK )) < 0 )
    { perror( SF ); exit( 1 ); }
    for( i = 1; i < argc; i++ )
    {
        if( strlen( argv[i] ) > PIPE_BUF - 2 )
            printf( "Too long: %s\n", argv[i] );
        else
        {
            make_msg( msgbuf, argv[i] );
            write( fd, msgbuf, PIPE_BUF );
        }
    }
    close( fd );
    return 0;
} /* end main */
```

send_msg.c cont'd

```
/* put input message into mb[] with '$' and padded with spaces */
void make_msg( char mb[], char input[])
{
    int i;
    for( i = 1; i < PIPE_BUF-1; i++ )
        mb[i] = ' ';
    mb[i] = '\0';
    i = 0;
    while( input[i] != '\0' )
    {
        mb[i] = input[i];
        i++;
    }
    mb[i] = '$';
} /* make_msg */
```


recv_msg.c

```
void print_msg( char mb[] );
int main( int argc, char *argv[] )
{
    int fd, I, done = 0;
    char msgbuf[PIPE_BUF];
    if(mkfifo(SF,0666) == -1)
        if(errno != EEXIST)
            { perror("receiver: mkfifo");
              exit( 1 ); }
    if( (fd = open( SF, O_RDWR )) < 0 )
        { perror( SF );
          exit( 1 ); }
    while( !done )
        {
            if( read( fd, msgbuf, PIPE_BUF ) < 0 )
                {
                    perror( "read" );
                    exit( 1 );
                }
            print_msg( msgbuf );
        }
    close( fd );
    return 0;
} /* end main */
```

recv_msg.c cont'd

```
/* print mb[] up to the '$' marker */
void print_msg( char mb[] )
{
    int i = 0;
    printf( "Msg: " );
    while( mb[i] != '$' )
    {
        putchar( mb[i] );
        i++;
    }
    putchar( '\n' );
} /* make_msg */
```

```
$ send_msg "Hello"
serv_fifo: No such file or directory
$ recv_msg &
[1] 8323
```

```
$ send_msg "Hello"
$ Msg: Hello
```

```
$ send_msg "Nice to see you"
Msg: Nice to see you
$ send_msg "This" "is" "four" "messages"
Msg: This
Msg: is
Msg: four
Msg: messages
$ kill -9 %1
[1] Killed          recv_msg
```