

UNIX System Programming

Socket Programming



Socket Programming

- What is a socket?
- Using sockets
 - Types (Protocols)
 - Associated functions
 - Styles

What is a socket

Socket API

- introduced in BSD4.1
UNIX, 1981
- Two types of sockets
 - connection-oriented
 - connectionless


socket

an interface (a “door”)
into which one
application process can
both send and
receive messages to/from
another (remote or
local) application process

Two essential types of sockets


- SOCK_STREAM

- a.k.a. TCP
- reliable delivery
- in-order guaranteed
- connection-oriented
- bidirectional




- SOCK_DGRAM

- a.k.a. UDP
- unreliable delivery
- no order guarantees
- no notion of “connection” – App indicates destination for each packet
- can send or receive




A Socket-eye view of the Internet



loki.trentu.ca
(192.197.151.116)



www.google.com
(74.125.226.145)




church.cse.ogi.edu
(129.95.50.2, 129.95.40.2)

- Each host machine has an IP address
- When a packet arrives at a host it contains information on which port it belongs to

Ports

- Each host has 65,536 ports
- Some ports are *reserved for specific apps*

- 20 & 21: FTP
- 22: SSH
- 23: Telnet
- 25: SMTP
- 53: DNS
- 80: HTTP
- 110: POP3
- 143: IMAP
- 443: HTTPS
- 465: SMTPS



- A socket provides an interface to send data to/from the network through a port

Addresses, Ports and Sockets

- Like apartments and mailboxes
 - You are the Application
 - Your home address is the IP address
 - Your mailbox is the port
 - The Post Office is the network
- Q: How do you choose which port to connects a socket?

Socket Creation in C: `socket`

- `int s = socket(domain, type, protocol);`
 - `s`: socket descriptor, an integer (like a file descriptor)
 - `domain`: integer, communication domain
 - `AF_INET` (IPv4 protocol) – typically used
 - `AF_INET6` (IPv6 protocol)
 - `AF_UNIX` or `AF_LOCAL` – intra-machine communication
 - `type`: communication type
 - `SOCK_STREAM`: reliable, 2-way, connection-based service
 - `SOCK_DGRAM`: unreliable, connectionless,
 - other values: need root permission, rarely used, or obsolete
 - `protocol`: specifies a particular protocol to be used with the socket. Normally only a single protocol exists to support a particular socket type within a given protocol family, in which case protocol can be specified as 0.
- NOTE: `socket` call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

Internet Addressing Data Structure

```
#include <netinet/in.h>

/* Internet address structure */
struct in_addr {
    u_long s_addr;          /* 32-bit IPv4 address */
};                          /* network byte ordered */

/* Socket address, Internet style. */
struct sockaddr_in {
    u_char sin_family;      /* Address Family */
    u_short sin_port;       /* UDP or TCP Port# */
                          /* network byte ordered */
    struct in_addr sin_addr; /* Internet Address */
    char sin_zero[8];       /* unused */
};
```

- `sin_family = AF_INET` selects Internet address family

Byte Ordering

```
union {  
    u_int32_t addr; /* 4 bytes address */  
    char c[4];  
} un;  
/* 128.2.194.95 */  
un.addr = 0x8002c25f;  
/* c[0] = ? */
```

- Big Endian


—Sun Solaris, PowerPC, ...

- Little Endian


—i386, alpha, ...

- Network byte order = Big Endian

c[0] c[1] c[2] c[3]



128	2	194	95
-----	---	-----	----



95	194	2	128
----	-----	---	-----

Address and port byte-ordering

- Address and port are stored as integers


—u_short sin_port; (16 bit)

—in_addr sin_addr; (32 bit)

```
struct in_addr {  
    u_long s_addr;  
};
```

□ Problem:


- different machines / OS's use different word orderings
 - little-endian: lower bytes first
 - big-endian: higher bytes first
- these machines may communicate with one another over the network



UNIX's byte-ordering funcs

- `u_long htonl(u_long x);`
- `u_short htons(u_short x);`
- `u_long ntohl(u_long x);`
- `u_short ntohs(u_short x);`

- ❑ On big-endian machines, these routines do nothing
- ❑ On little-endian machines, they reverse the byte order



- ❑ Same code would have worked regardless of endianness of the two machines


Byte Ordering Functions

- Converts between **host byte order** and **network byte order**
 - ‘h’ = host byte order
 - ‘n’ = network byte order
 - ‘l’ = long (4 bytes), converts IP addresses
 - ‘s’ = short (2 bytes), converts port numbers

```
#include <netinet/in.h>

unsigned long int htonl(unsigned long int hostlong);
unsigned short int htons(unsigned short int
hostshort);
unsigned long int ntohl(unsigned long int netlong);
unsigned short int ntohs(unsigned short int
netshort);
```

TCP Server



- For example: web server
- **What does a *web server* need to do so that a *web client* can connect to it?**

Socket I/O: socket()

- Since web traffic uses TCP, the web server must create a socket of type SOCK_STREAM (connection-oriented)

```
int socketd;                /* socket descriptor */

if((socketd=socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- *socket* returns an integer (socket descriptor)
 - socket* < 0 indicates that an error occurred
- **AF_INET** associates a socket with the Internet protocol family
- **SOCK_STREAM** selects the TCP protocol
- In some Unixes, need to compile with **-lsocket -lnsl** to link in socket libraries

Socket I/O: bind()

- A *socket* can be bound to a *port*

```
int socketd;                                /* socket descriptor */
struct sockaddr_in srv;                    /* used by bind() */

/* create the socket */

srv.sin_family = AF_INET; /* use the Internet addr family */
srv.sin_port = htons(80); /* bind socket 'socketd' to port 80*/

/* bind: a client may connect to any of my addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);
/* INADDR_ANY - refers to local machine address */

if(bind(socketd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("bind"); exit(1);
}
```

- Still not quite ready to communicate with a client...

Socket I/O: listen()

- *listen* indicates that the server will accept a connection

```
int socketd;          /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* 1) create the socket */
/* 2) bind the socket to a port */

if(listen(socketd, 5) < 0) {
    perror("listen");
    exit(1);
}
```

- **Still not quite ready to communicate with a client...**

Socket I/O: accept()

- *accept* blocks waiting for a connection

```
int socketd;                /* socket descriptor */
struct sockaddr_in srv;     /* used by bind() */
struct sockaddr_in cli;     /* used by accept() */
int newfd;                  /* returned by accept() */
int cli_len = sizeof(cli);  /* used by accept() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */

newfd = accept(socketd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept");    exit(1);
}
```

- *accept* returns a new socket (*newfd*) with the same properties as the original socket (socketd)
—*newfd* < 0 indicates that an error occurred

Socket I/O: accept() continued...

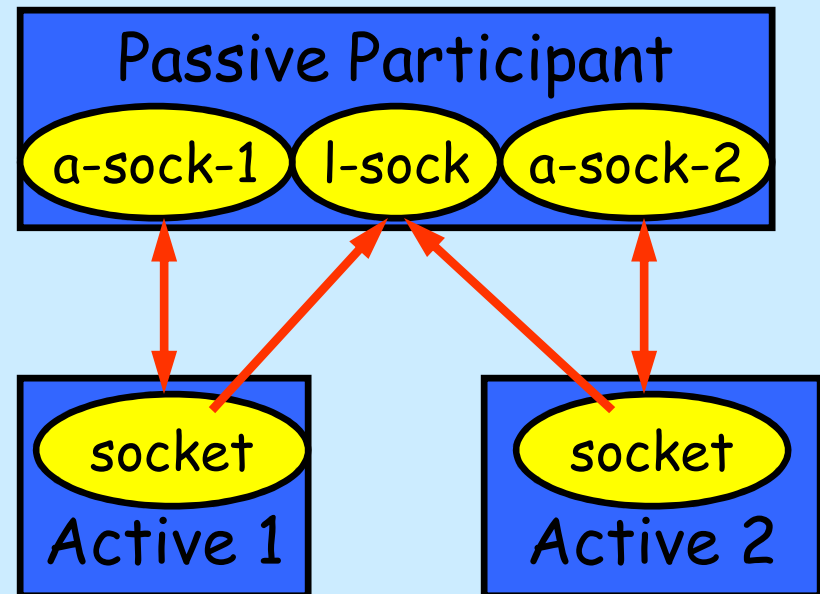
```
struct sockaddr_in cli;          /* used by accept() */
int newfd;                      /* returned by accept() */
int cli_len = sizeof(cli);      /* used by accept() */

newfd = accept(socketd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept");
    exit(1);
}
```

- How does the server know which client it is?
 - cli.sin_addr.s_addr** contains the client's *IP address*
 - cli.sin_port** contains the client's *port number*
- Now the server can exchange data with the client by using *read* and *write* on the descriptor *newfd*.
- Why does *accept* need to return a new descriptor?

Connection setup

- Passive participant - server
 - step 1: listen (for incoming requests)
 - step 3: accept (a request)
 - step 4: data transfer
 - Active participant - clients
 - step 2: request & establish connection
 - step 4: data transfer
- The accepted connection is on a new socket
- The old socket continues to listen for other active participants



Socket I/O: read()

- *read* can be used with a socket
- *read* **blocks** waiting for data from the client but does not guarantee that `sizeof(buf)` is read


```
int sockfd;                /* socket descriptor */
char buf[512];             /* used by read() */
int nbytes;                /* used by read() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */
/* 4) accept the incoming connection */

if((nbytes = read(sockfd, buf, sizeof(buf))) < 0) {
    perror("read"); exit(1);
}
```

TCP Client

- For example: web client
- **How does a *web client* connect to a *web server*?**



Dealing with IP Addresses

- IP Addresses are commonly written as strings (“128.192.35.50”), but programs deal with IP addresses as integers.

Converting strings to numerical address:

```
struct sockaddr_in srv;  
  
srv.sin_addr.s_addr = inet_addr("128.192.35.50") ;  
if(srv.sin_addr.s_addr == (in_addr_t) -1) {  
    fprintf(stderr, "inet_addr failed!\n"); exit(1);  
}
```

Converting a numerical address to a string:

```
struct sockaddr_in srv;  
char *t = inet_ntoa(srv.sin_addr) ;  
if(t == 0) {  
    fprintf(stderr, "inet_ntoa failed!\n"); exit(1);  
}
```

Translating Names to Addresses

- Gethostbyname provides interface to DNS
- Additional useful calls
 - Gethostbyaddr – returns hostent given sockaddr_in
 - Getservbyname
 - Used to get service description (typically port number)
 - Returns servent based on name

```
#include <netdb.h>

struct hostent *hp; /*ptr to host info for remote*/
struct sockaddr_in peeraddr;
char *name = "odin.trentu.ca";

peeraddr.sin_family = AF_INET;
hp = gethostbyname(name)
peeraddr.sin_addr.s_addr = ((struct in_addr*) (hp->h_addr))->s_addr;
```


Socket I/O: connect()

- *connect* allows a client to connect to a server...

```
int socketd;                                /* socket descriptor */
struct sockaddr_in srv;                     /* used by connect() */

/* create the socket */

/* connect: use the Internet address family */
srv.sin_family = AF_INET;

/* connect: socket 'socketd' to port 80 */
srv.sin_port = htons(80);

/* connect: connect to IP Address "192.197.151.70" */
srv.sin_addr.s_addr = inet_addr("192.197.151.70");

if(connect(socketd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("connect"); exit(1);
}
```

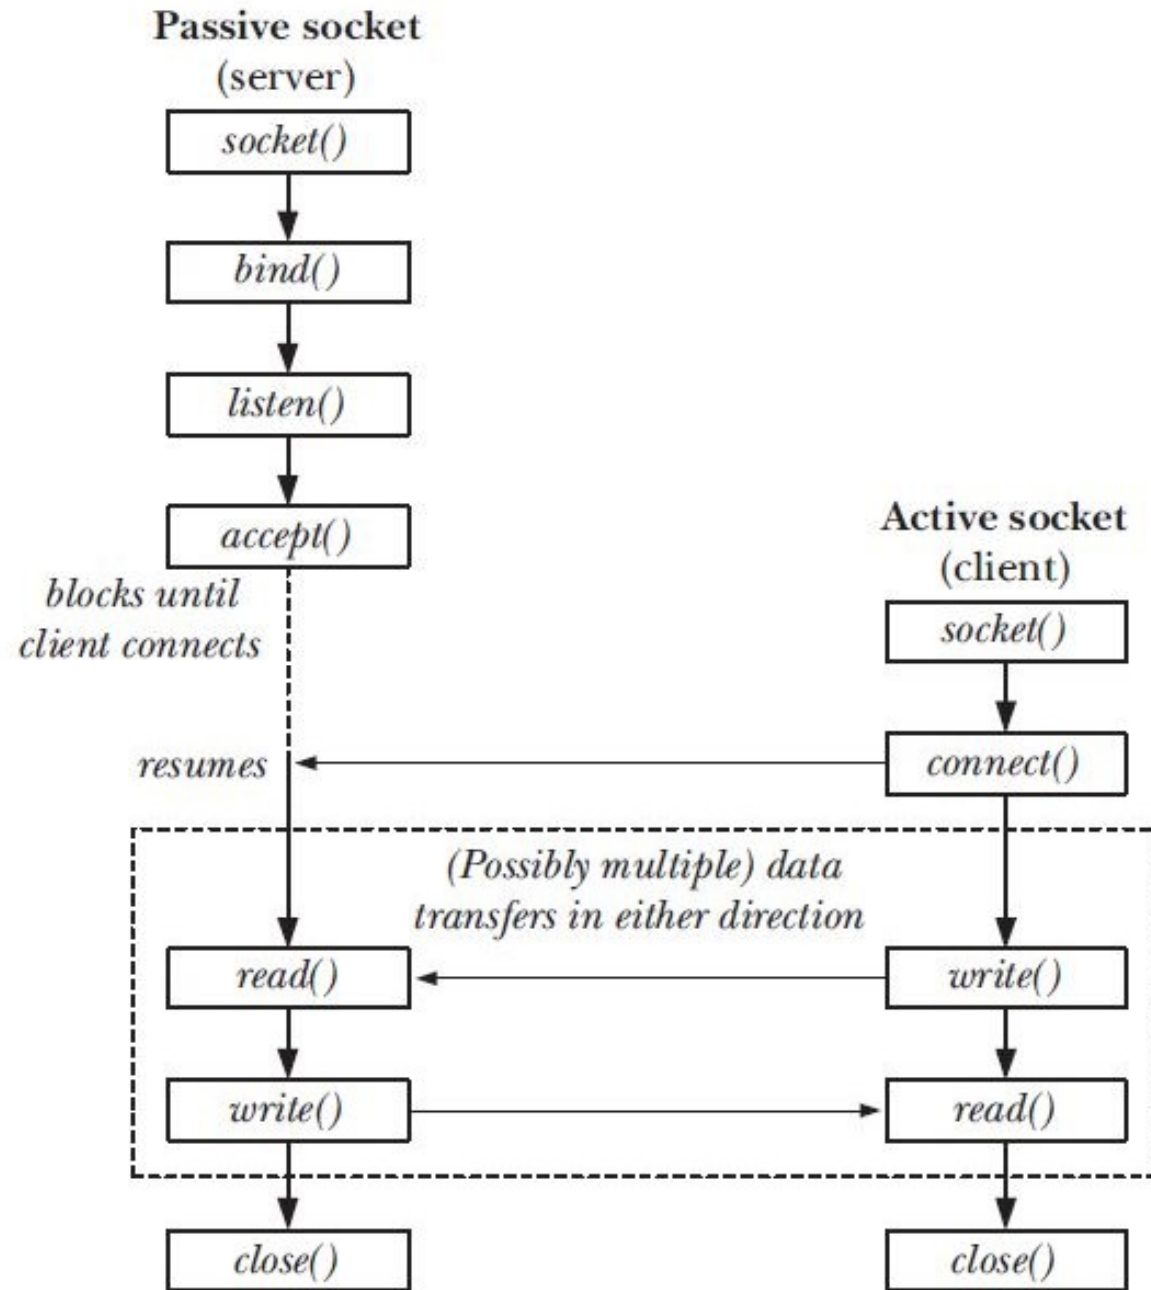
Socket I/O: write()

- *write* can be used with a socket

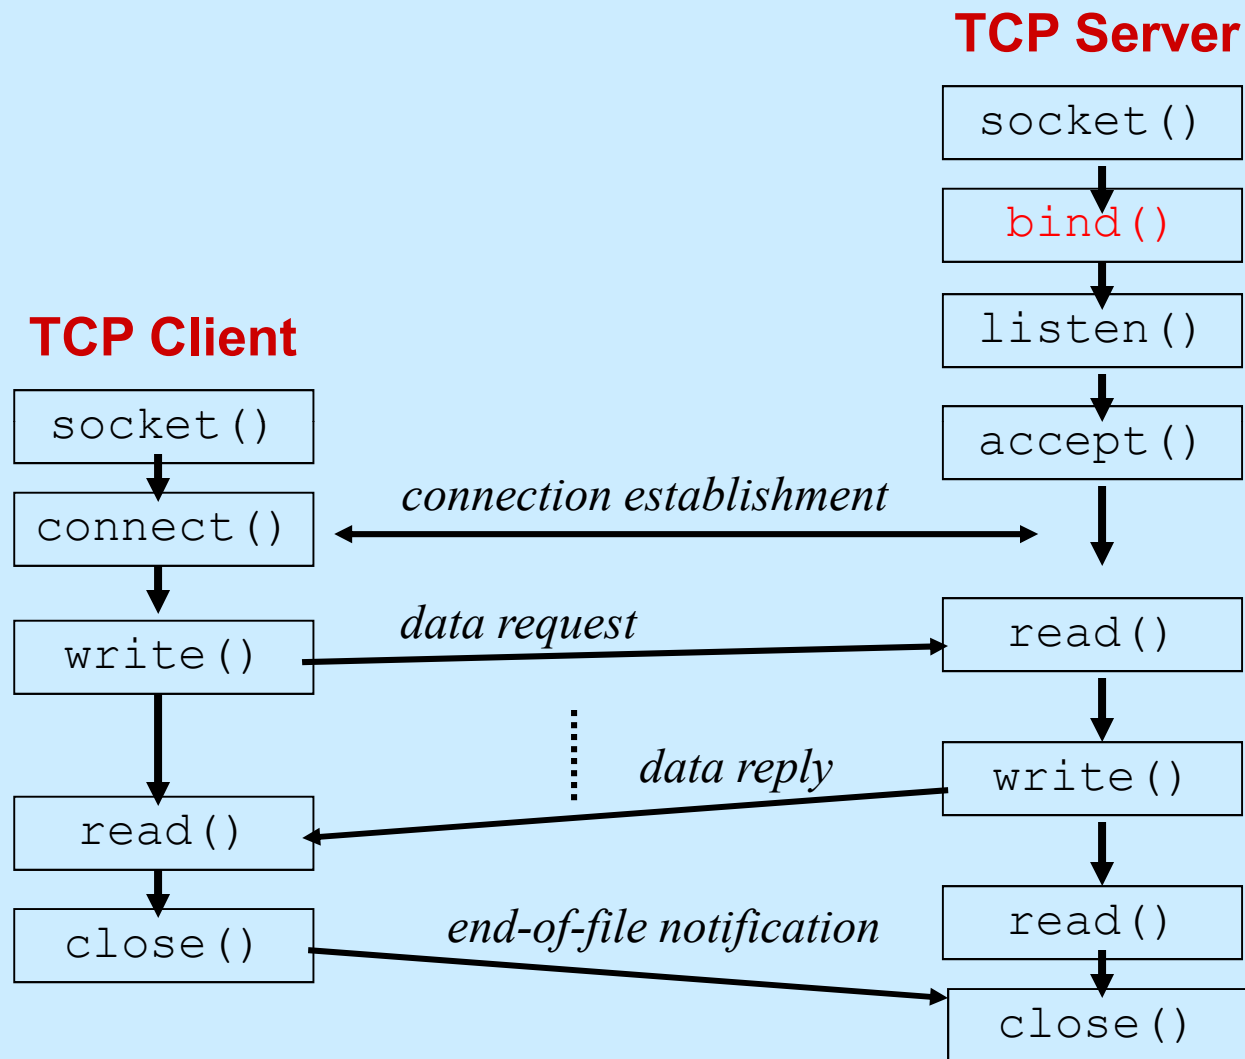
```
int socketd;                /* socket descriptor */
struct sockaddr_in srv;     /* used by connect() */
char buf[512];             /* used by write() */
int nbytes;                /* used by write() */

/* 1) create the socket */
/* 2) connect() to the server */

/* Example: A client could "write" a request to a server */
if((nbytes = write(socketd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```



Review: TCP Client-Server Interaction



The *bind* function

- associates and (can exclusively) reserves a port for use by the socket
- `int status = bind(sockid, &addrport, size);`
 - `status`: error status, = -1 if bind failed
 - `sockid`: integer, socket descriptor
 - `addrport`: struct `sockaddr`, the (IP) address and port of the machine (address usually set to `INADDR_ANY` – chooses a local address)
 - `size`: the size (in bytes) of the `addrport` structure
- `bind` can be skipped for both types of sockets.

Skipping the *bind*

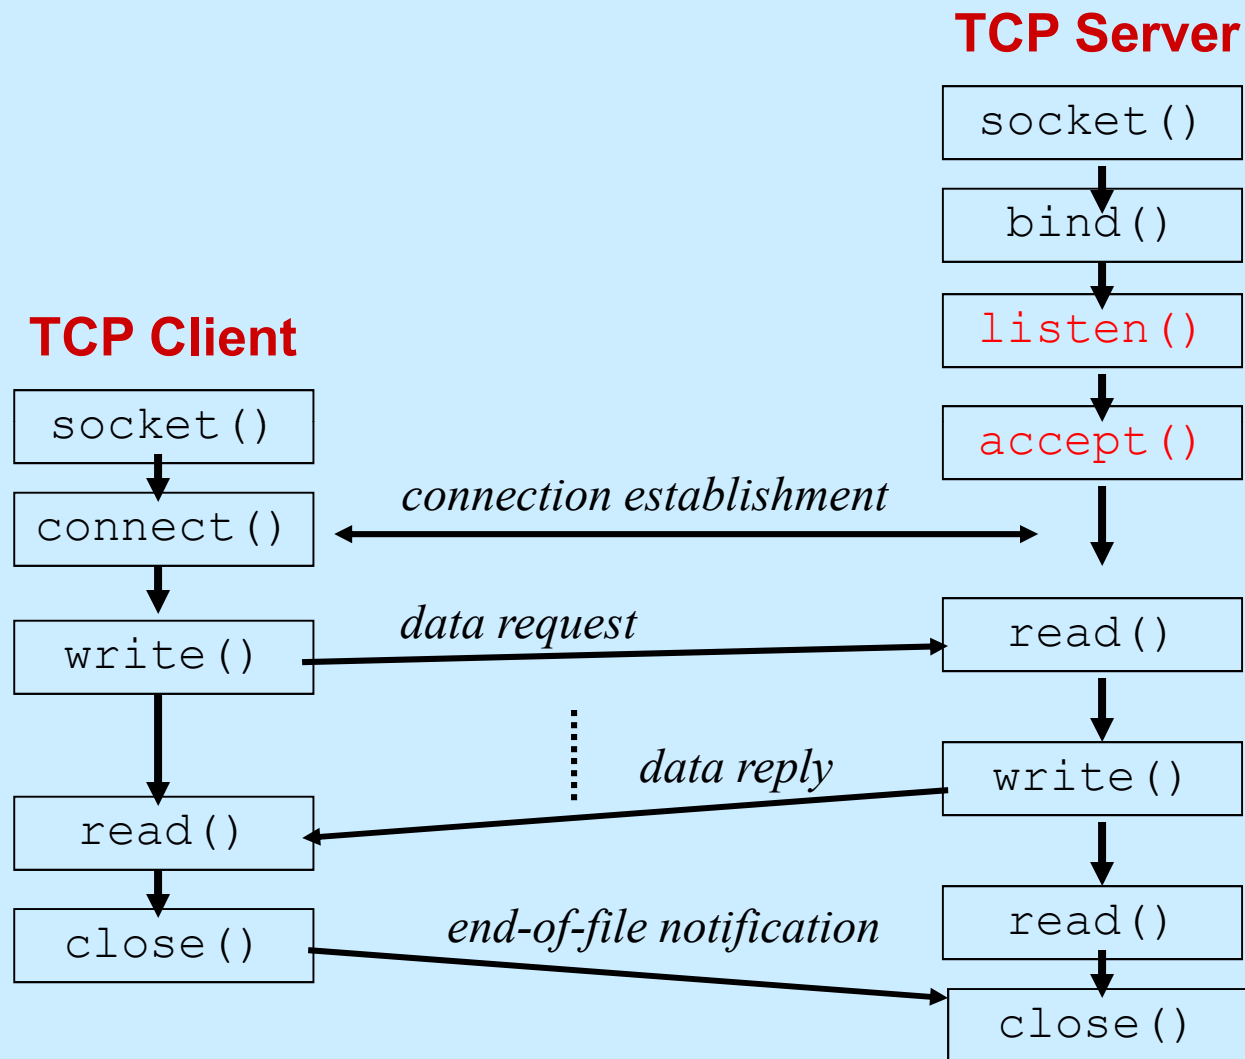
- **SOCK_DGRAM:**

- if only sending, no need to bind. The OS finds a port each time the socket sends a packet
- if receiving, need to bind

- **SOCK_STREAM:**

- destination determined during connection setup
- don't need to know port sending from (during connection setup, receiving end is informed of port)

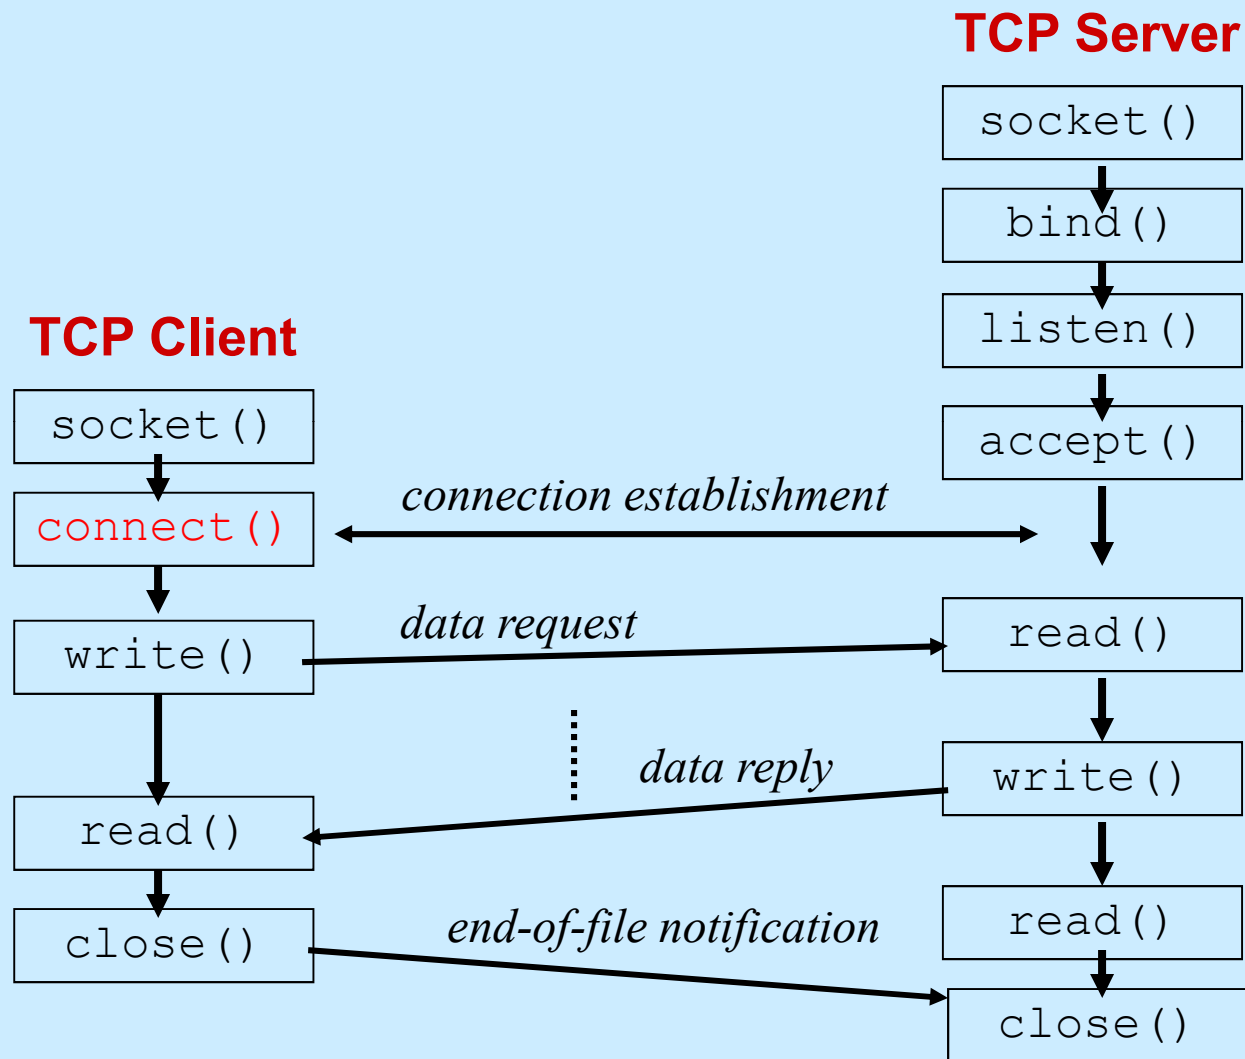
Review: TCP Client-Server Interaction



Connection setup: *listen* & *accept*

- Called by passive participant
- `int status = listen(sock, queuelen);`
 - `status`: 0 if listening, -1 if error
 - `sock`: integer, socket descriptor
 - `queuelen`: integer, # of active participants that can “wait” for a connection
 - `listen` is **non-blocking**: returns immediately
- `int s = accept(sock, &name, &namelen);`
 - `s`: integer, the new socket (used for data-transfer)
 - `sock`: integer, the orig. socket (being listened on)
 - `name`: struct `sockaddr`, address of the active participant
 - `namelen`: `sizeof(name)`: value/result parameter
 - must be set appropriately before call
 - adjusted by OS upon return
 - `accept` is **blocking**: waits for connection before returning

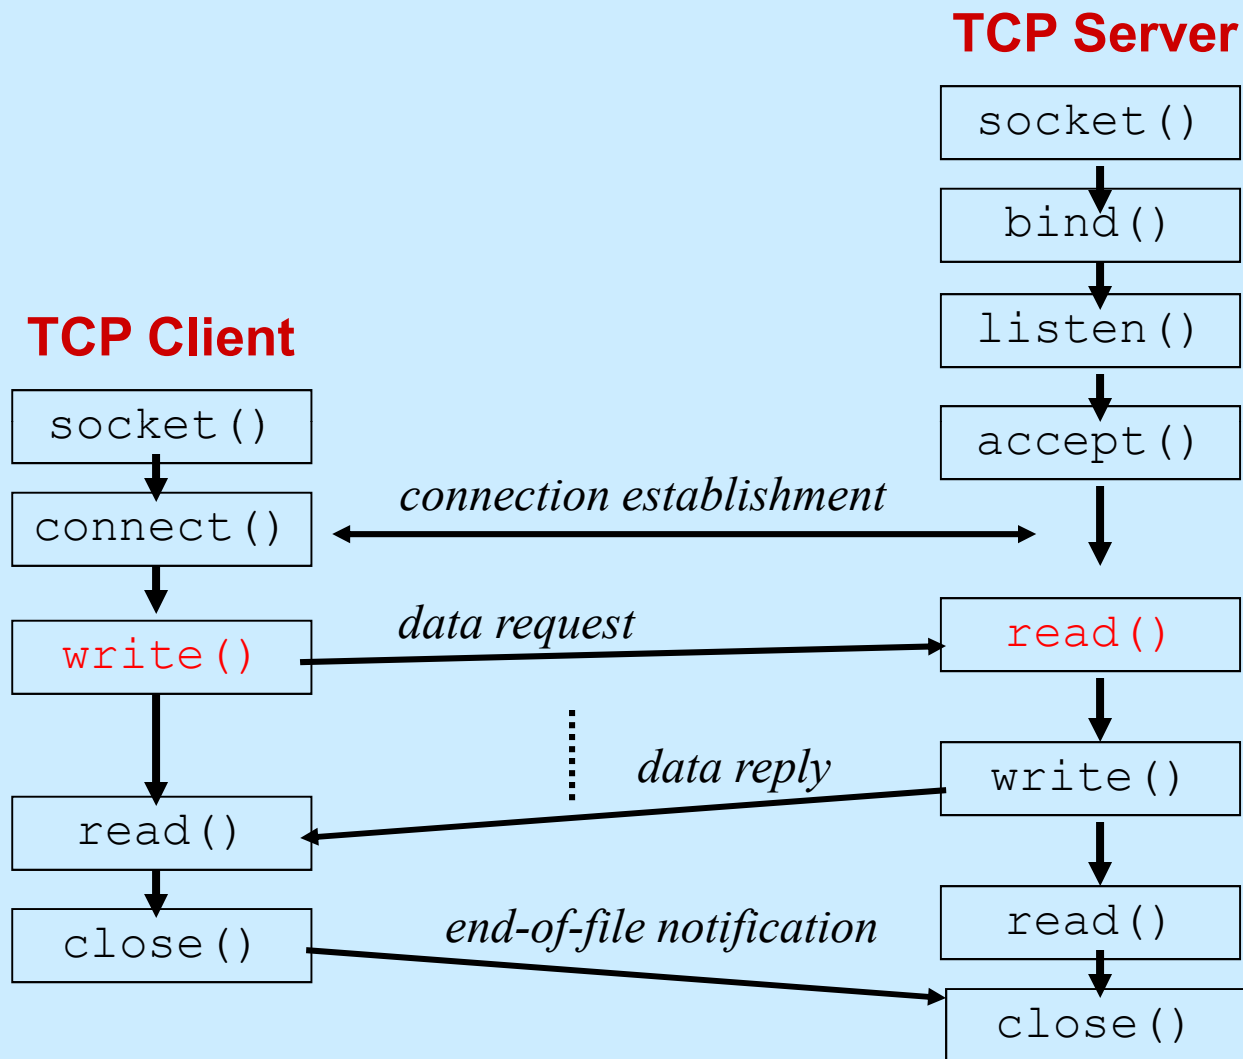
Review: TCP Client-Server Interaction



connect call

- `int status = connect(sock, &name, namelen);`
 - status: 0 if successful connect, -1 otherwise
 - sock: integer, socket to be used in connection
 - name: struct `sockaddr`: address of passive participant
 - namelen: integer, `sizeof(name)`
- `connect` is **blocking**

Review: TCP Client-Server Interaction



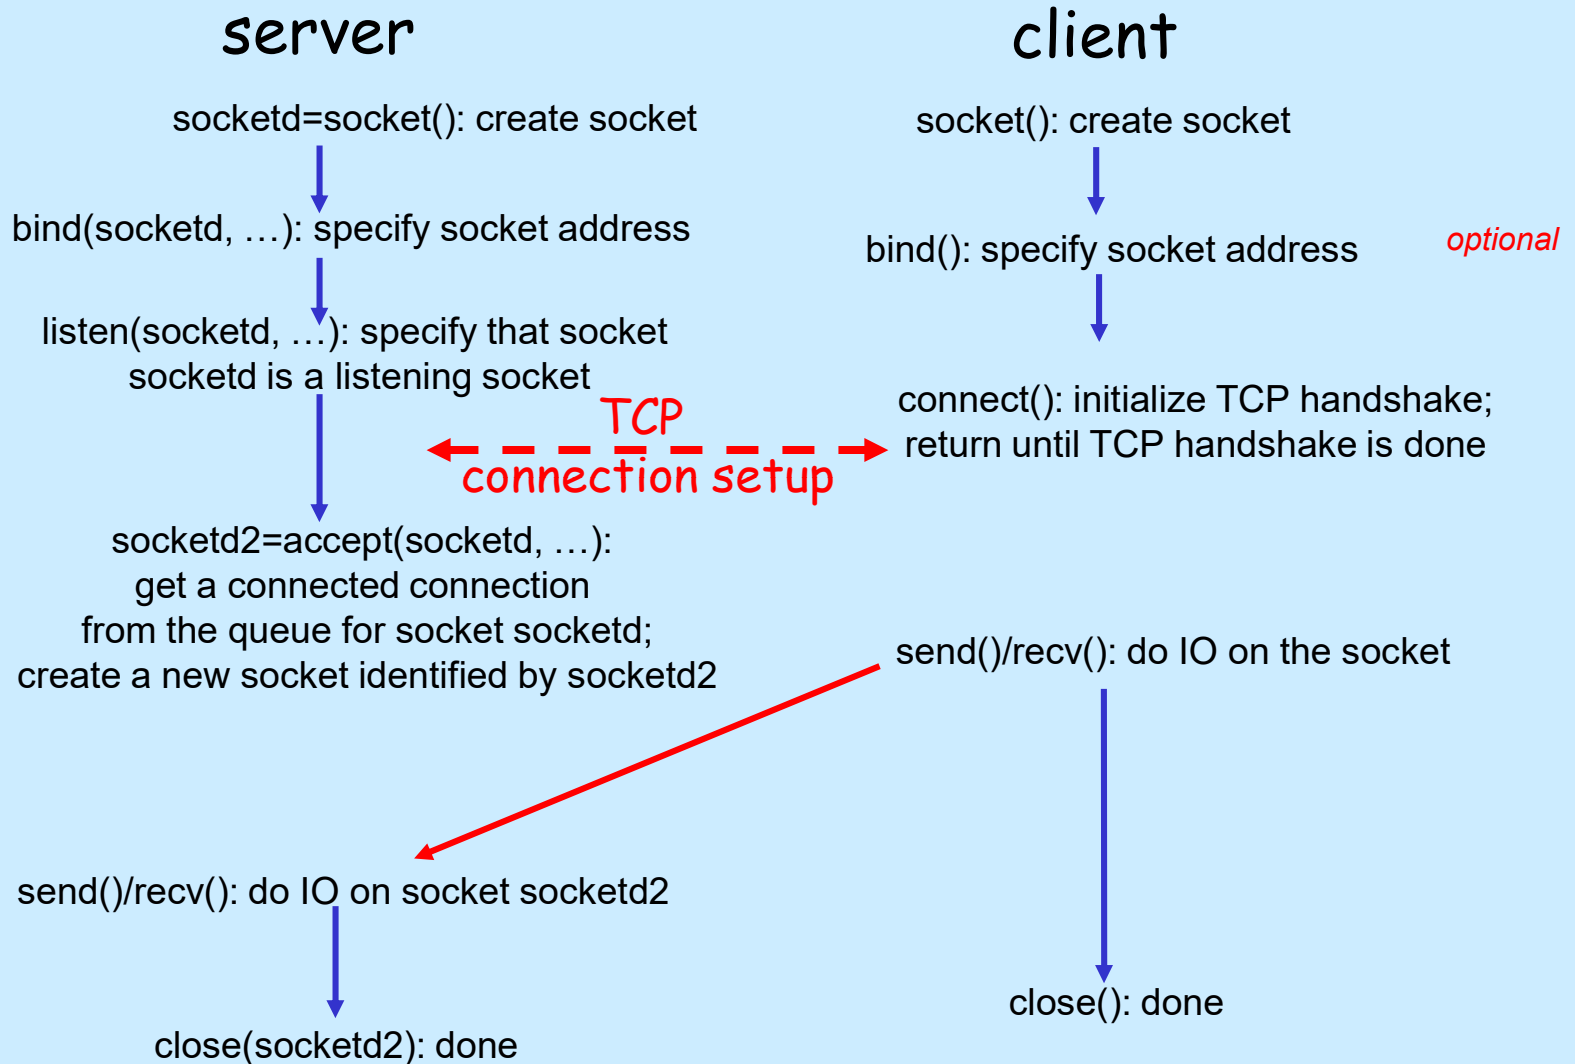
Sending / Receiving Data

- With a connection (**SOCK_STREAM**):
 - Use `send()/recv()` instead of `read()/write()`
 - `int count = send(sock, &buf, len, flags);`
 - `count`: # bytes transmitted (-1 if error)
 - `buf`: `char[]`, buffer to be transmitted
 - `len`: integer, length of buffer (in bytes) to transmit
 - `flags`: integer, special options, usually just 0
 - `int count = recv(sock, &buf, len, flags);`
 - `count`: # bytes received (-1 if error)
 - `buf`: `void[]`, stores received bytes
 - `len`: # bytes received
 - `flags`: integer, special options, usually just 0
 - Calls are **blocking** [returns only after data is sent (to socket buf) / received]

close

- When finished using a socket, the socket should be closed:
- `status = close(s);`
 - **status**: 0 if successful, -1 if error
 - **s**: the file descriptor (socket being closed)
- Closing a socket
 - closes a connection (for `SOCK_STREAM`)
 - frees up the port used by the socket

Connection-oriented: Big Picture



Connection Example - server

```
#include <sys/socket.h> /* for socket(), connect(), send(), and recv() */
#include <netinet/in.h>
#include <signal.h>
#include <ctype.h>
void catcher(int sig);
int newfd;
main()
{
    int socketd;                /* socket descriptor */
    struct sockaddr_in srv;     /* used by bind() */
    struct sockaddr_in cli;     /* used by accept() */
    int cli_len = sizeof(cli);  /* used by accept() */
    int not_done = 1;
    char c;

    signal(SIGPIPE, catcher);

    /* 1) create the socket */
    if((socketd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        perror("socket call failed");
        exit(1); }
}
```

Connection Example - server

```
/* 2) bind the socket to a port */  
    srv.sin_family = AF_INET; /* use the Internet addr family */  
    srv.sin_port = htons(1700); /* bind socket 'socketd' to port 1700 */  
  
    /* bind: a client may connect to any of my addresses */  
    srv.sin_addr.s_addr = htonl(INADDR_ANY);  
                                /* INADDR_ANY - refers to local machine address */  
  
    if(bind(socketd, (struct sockaddr *) &srv, sizeof(srv)) < 0) {  
        perror("bind call failed");  
        exit(1);  
    }
```


Connection example – server con't

```
/* 3) listen on the socket */  
if(listen(socketd, 5) < 0) {  
    perror("listen call failed");  
    exit(1);  
}  
  
/* loop looking for messages */  
while (not_done)  
    {  
        /* 4) accept the incoming connection */  
        newfd = accept(socketd, (struct sockaddr*) &cli, &cli_len);  
        if(newfd < 0) {  
            perror("accept call failed");  
            not_done = 0;  
        } // endif
```

Connection example – server con't

```
/* spawn a child to deal with this connection */  
if ( fork() == 0)  
{  
    while(recv(newfd, &c, 1, 0) > 0) /* could use read as well */  
    {  
        c = toupper(c);  
        send(newfd, &c, 1, 0); /* could use write as well */  
    }  
    /* when client is no longer sending, close socket and child */  
    close(newfd);  
    exit(0);  
}  
    else /* parent */ { close(newfd); }  
} // While not done
```

```
/* signal handler in case socket becomes disconnected */  
void catcher(int sig)  
{  
    signal(SIGPIPE, catcher);  
    close(newfd);  
    exit(0);  
}
```

Connection Example - client

```
#include <sys/socket.h> /* for socket(), connect(), send(), and recv() */
#include <arpa/inet.h> /* for sockaddr_in and inet_addr() */
#include <netinet/in.h>
#include <signal.h>
main()
{
    int socketd;                /* socket descriptor */
    struct sockaddr_in srv;      /* used by connect() */
    char c, rc;
    int more_data = 1;

    /* 1) create the socket */

    if((socketd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        perror("socket call failed");
        exit(1);
    }
```

Connection Example - client

```
/* 2) connect() to the server */
```

```
srv.sin_family = AF_INET; /* connect: use the Internet address family */  
srv.sin_port = htons(1700); /* connect: socket 'socketd' to port 1700 */  
srv.sin_addr.s_addr = inet_addr("192.197.151.70"); /* connect: connect to IP Address  
                                '192.197.151.70' - odin */
```

```
if(connect(socketd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {  
    perror("connect call failed"); exit(1);  
}
```

Connection example – client con't

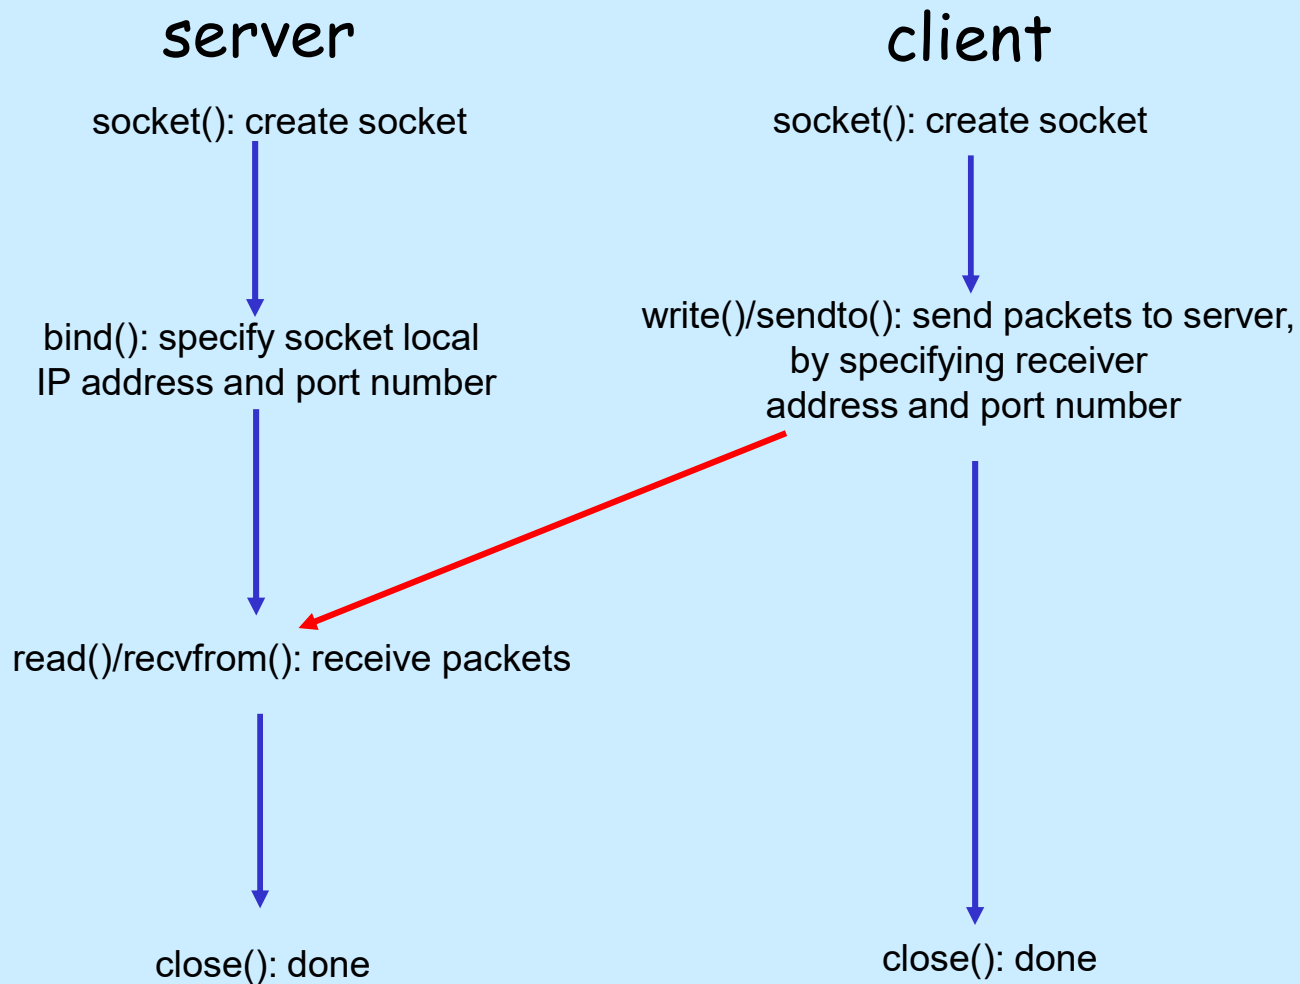
```
/* send and receive information with the server */
while (more_data)
{
    if(c != '\n') /* ignore the enter at the end of the input */
        printf("Input a lower case letter (or 0 to stop) => ");
    c = getchar();

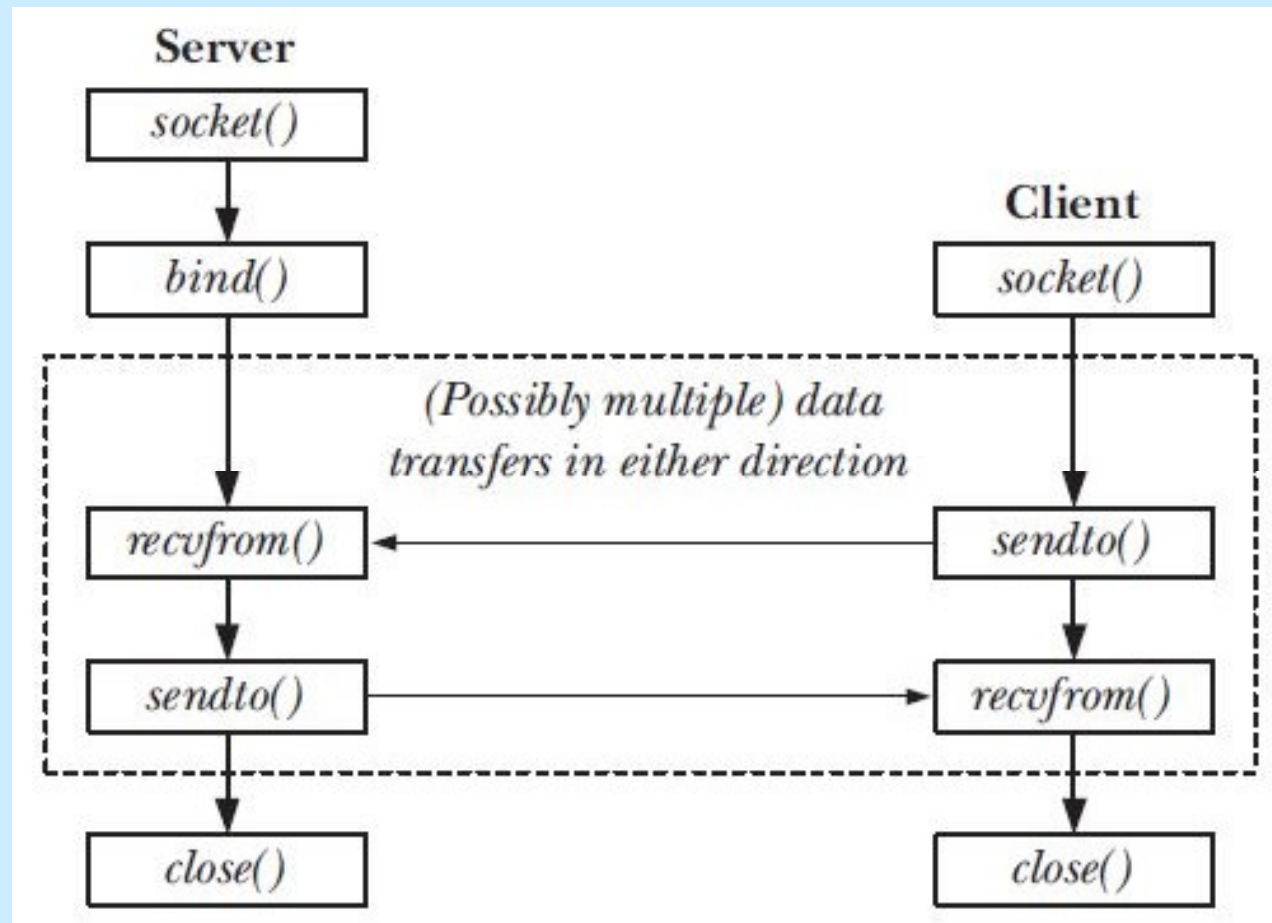
    if(c != '0')
    {
        if(c != '\n') /* ignore the enter for processing */
        {
            send(socketd, &c, 1, 0); /* could use write as well */
            if(recv(socketd, &rc, 1, 0) > 0) /* could use read as well */
                printf("%c\n", rc);
            else {
                printf("Server has died\n");
                close(socketd);
                exit(1);
            } // recv
        } // if not NewLine
    } // if not a ZERO
    else
        more_data = 0;
} // end of while
exit(0);
}
```

Connection Example


- To compile use
 - gcc -o server server.c // some versions need -lsocket -lnsl
 - gcc -o client client.c // some versions need -lsocket -lnsl
- To run, start the server first in the background and then the client(s)
 - \$ server &
 - \$ client
 - Input a lower case letter (or 0 to stop) => r
 - R
 - Input a lower case letter (or 0 to stop) => i
 - I
 - Input a lower case letter (or 0 to stop) => c
 - C
 - Input a lower case letter (or 0 to stop) => h
 - H
 - Input a lower case letter (or 0 to stop) => 0
 - \$ kill %1
 - [1] Terminated server

Connectionless: Big Picture





UDP Server Example



- For example: NTP daemon
- **What does a *UDP* server need to do so that a *UDP* client can connect to it?**

Socket I/O: socket()

- The UDP server must create a **datagram** socket...

```
int socketd;          /* socket descriptor */

if((socketd = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- *socket* returns an integer (**socket descriptor**)
—*socketd* < 0 indicates that an error occurred
- AF_INET: associates a socket with the Internet protocol family
- **SOCK_DGRAM**: selects the UDP protocol

Socket I/O: bind()

- A *socket* can be bound to a *port*

```
int socketd;                                /* socket descriptor */
struct sockaddr_in srv;                     /* used by bind() */

/* create the socket */

/* bind: use the Internet address family */
srv.sin_family = AF_INET;

/* bind: socket 'socketd' to port 80*/
srv.sin_port = htons(80);

/* bind: a client may connect to any of my addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);

if(bind(socketd, (struct sockaddr*) &srv, sizeof(srv)) < 0)
{
    perror("bind"); exit(1);
}
```

- Now the UDP server is ready to accept packets...

Socket I/O: recvfrom()

- *read* does not provide the client's address to the UDP server

```
int socketd;                /* socket descriptor */
struct sockaddr_in srv;     /* used by bind() */
struct sockaddr_in cli;     /* used by recvfrom() */
char buf[512];              /* used by recvfrom() */
int cli_len = sizeof(cli);  /* used by recvfrom() */
int nbytes;                 /* used by recvfrom() */

/* 1) create the socket */
/* 2) bind to the socket */

nbytes = recvfrom(socketd, buf, sizeof(buf), 0,
                  (struct sockaddr*) &cli, &cli_len);
if(nbytes < 0) {
    perror("recvfrom"); exit(1);
}
```

Socket I/O: `recvfrom()` continued...


```
nbytes = recvfrom(socketd, buf, sizeof(buf), 0 /* flags */,  
                  (struct sockaddr*) cli, &cli_len);
```

- The actions performed by *recvfrom*

- returns the number of bytes read (*nbytes*)
- copies *nbytes* of data into *buf*
- returns the address of the client (*cli*)
- returns the length of *cli* (*cli_len*)
- don't worry about flags

UDP Client Example

- How does a *UDP client* communicate with a *UDP server*?



Socket I/O: sendto()

- *write* is not allowed
- Notice that the UDP client does not *bind* a port number
 - a port number is **dynamically assigned** when the first *sendto* is called

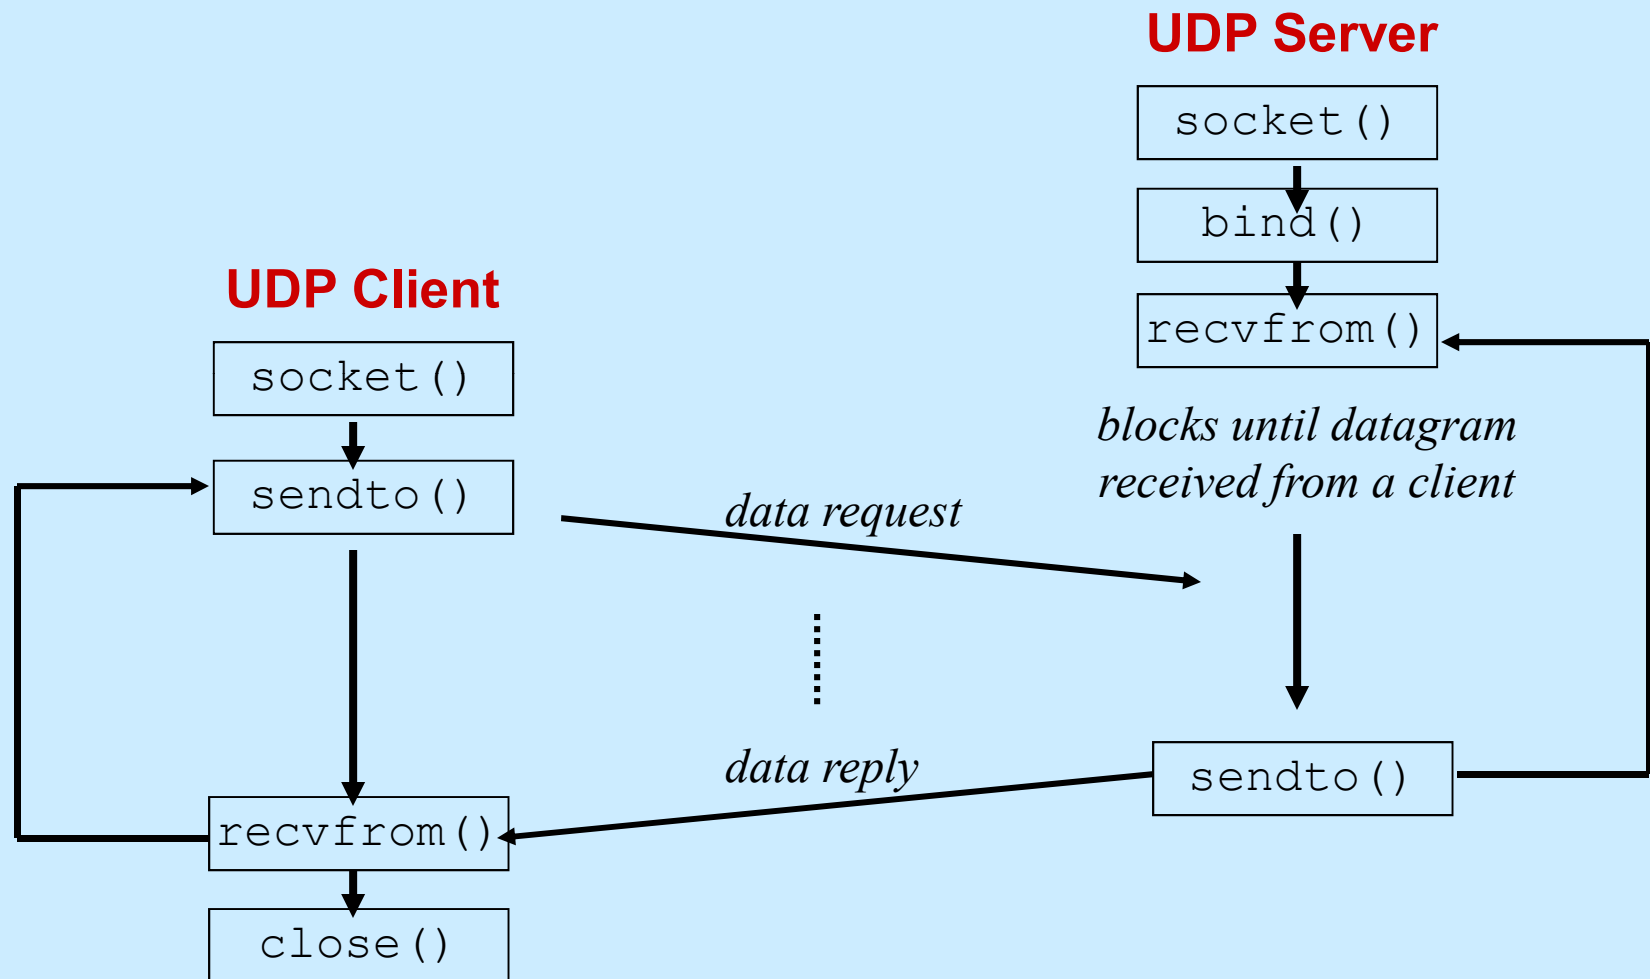
```
int socketd;                /* socket descriptor */
struct sockaddr_in srv;     /* used by sendto() */

/* 1) create the socket */

/* sendto: send data to IP Address "192.197.151.70" port 80 */
srv.sin_family = AF_INET;
srv.sin_port = htons(80);
srv.sin_addr.s_addr = inet_addr("192.197.151.70");

nbytes = sendto(socketd, buf, sizeof(buf), 0 /* flags */,
                (struct sockaddr*) &srv, sizeof(srv));
if(nbytes < 0) {
    perror("sendto");    exit(1);
}
```

Review: UDP Client-Server Interaction



Connectionless example - server

```
#include <sys/socket.h> /* for socket(), connect(), send(), and recv() */
#include <netinet/in.h>
#include <signal.h>
#include <ctype.h>
main()
{
    int socketd;                /* socket descriptor */
    struct sockaddr_in srv;     /* used by bind() */
    struct sockaddr_in cli;     /* used by sendto(), recvfrom() */
    int cli_len = sizeof(cli);  /* used by sendto(), recvfrom() */
    int not_done = 1;
    char c;
    /* 1) create the socket */
    if((socketd = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
        perror("socket call failed");
        exit(1);
    }
}
```

Connectionless example - server

```
/* 2) bind the socket to a port */  
srv.sin_family = AF_INET; /* use the Internet addr family */  
  
srv.sin_port = htons(1700); /* bind socket 'socketd' to port 1700 */  
  
/* bind: a client may connect to any of my addresses */  
srv.sin_addr.s_addr = htonl(INADDR_ANY);  
/* INADDR_ANY - refers to local machine address */  
  
if(bind(socketd, (struct sockaddr *) &srv, sizeof(srv)) < 0) {  
    perror("bind call failed");  
    exit(1);  
}
```

Connectionless – server con't

```
/* loop looking for messages */
while(not_done)
{
    /* 3) receive a message */
    if(recvfrom(socketd, &c, 1, 0, (struct sockaddr *) &cli, &cli_len) < 0)
    {
        perror("Server: receiving");
        not_done = 0;
    }

    c = toupper(c);

    /* send the message back to where it came from */
    if(sendto(socketd, &c, 1, 0, (struct sockaddr *) &cli, cli_len) < 0)
    {
        perror("Server: sending");
        not_done = 0;
    }
} // while not done
exit(0);
}
```

Connectionless example - client

```
#include <sys/socket.h> /* for socket(), connect(), send(), and recv() */
#include <arpa/inet.h> /* for sockaddr_in and inet_addr() */
#include <netinet/in.h>
#include <signal.h>
main()
{
    int socketd;                /* socket descriptor */
    struct sockaddr_in srv;      /* used by sendto() */
    int srv_len = sizeof(srv);
    char c;
    int more_data = 1;
    /* 1) create the socket */
    if((socketd = socket(AF_INET, SOCK_DGRAM, 0)) < 0) {
        perror("socket call failed");
        exit(1);
    }
}
```

Connectionless example - client

```
/* set up the server information for the sendto */  
srv.sin_family = AF_INET;  
/* port number */  
srv.sin_port = htons(1700);  
/* IP Address '192.197.151.70' - odin */  
srv.sin_addr.s_addr = inet_addr("192.197.151.70");
```

Connectionless – client con't


```
/* send and receive information with the server */
while (more_data) {
    if(c != '\n') /* ignore the enter at the end of the input */
        printf("Input a lower case letter (or 0 to stop) => ");
    c = getchar();
    if(c != '0') {
        if(c != '\n') { /* ignore the enter for processing */

            if (sendto(socketd, &c, 1, 0, (struct sockaddr *) &srv, srv_len) < 0) {
                perror("Client: sending");
                exit(1);
            }
            if (recvfrom(socketd, &c, 1, 0, (struct sockaddr *) &srv, &srv_len) < 0) {
                perror("Client: receiving");
                exit(1);
            }
            printf("%c\n", c);
        } // c!= new line
    } // c !=0
    else
        more_data = 0;
}
exit(0);
}
```

Connectionless Example

- To compile use
 - gcc -o server1 server1.c // some versions need -lsocket -lnsl
 - gcc -o client1 client1.c // some versions need -lsocket -lnsl
- To run, start the server first in the background and then the client(s)
 - \$ server1 &
 - \$ client1
 - Input a lower case letter (or 0 to stop) => r
 - R
 - Input a lower case letter (or 0 to stop) => i
 - I
 - Input a lower case letter (or 0 to stop) => c
 - C
 - Input a lower case letter (or 0 to stop) => h
 - H
 - Input a lower case letter (or 0 to stop) => 0
 - \$ kill %1
 - [1] Terminated server1

The UDP Server



- How can the *UDP* server service multiple ports simultaneously?

UDP Server: Servicing Two Ports

```
int s1;                /* socket descriptor 1 */
int s2;                /* socket descriptor 2 */
int not_done = 1;

/* 1) create socket s1 */
/* 2) create socket s2 */
/* 3) bind s1 to port 2000 */
/* 4) bind s2 to port 3000 */

while(not_done) {
    recvfrom(s1, buf, sizeof(buf), ...);
    /* process buf */

    recvfrom(s2, buf, sizeof(buf), ...);
    /* process buf */
}
```

- **What problems does this code have?**

Dealing with blocking calls

- Many of the functions we saw block until a certain event
 - accept: until a connection comes in
 - connect: until the connection is established
 - recv, recvfrom: until a packet (of data) is received
 - send, sendto: until data is pushed into socket's buffer
- For simple programs, blocking is convenient
- What about more complex programs?
 - multiple connections
 - simultaneous sends and receives
 - simultaneously doing non-networking processing

Dealing w/ blocking (cont'd)

- Options:
 - create multi-process or multi-threaded code
 - turn off the blocking feature (e.g., using the `fcntl` file-descriptor control function)
 - use the **select** function call.
- What does **select()** do?
 - can be permanent blocking, time-limited blocking or non-blocking
 - input: a set of file-descriptors
 - output: info on the file-descriptors' status
 - i.e., can identify sockets that are “ready for use”: calls involving that socket will return immediately

select function call

- `int status = select(nfds, &readfds, &writefds, &exceptfds, &timeout);`
 - `status`: # of ready objects, -1 if error
 - `nfds`: 1 + largest file descriptor to check
 - `readfds`: list of descriptors to check if read-ready
 - `writefds`: list of descriptors to check if write-ready
 - `exceptfds`: list of descriptors to check if an exception is registered
 - `timeout`: time after which select returns, even if nothing ready - can be 0 or ∞
(point timeout parameter to `NULL` for ∞)

To be used with select:

- Recall **select** uses a structure, **struct fd_set**
 - it is just a bit-vector
 - if bit i is set in [readfds, writefds, exceptfds], select will check if file descriptor (i.e. socket) i is ready for [reading, writing, exception]
- Before calling **select**:
 - **FD_ZERO(&fdvar)**: clears the structure
 - **FD_SET(i , &fdvar)**: to check file desc. i
- After calling **select**:
 - **int FD_ISSET(i , &fdvar)**: boolean returns **TRUE** iff i is “ready”

Socket I/O: select()

```
int select(int maxfds, fd_set *readfds, fd_set *writelfds,  
           fd_set *exceptfds, struct timeval *timeout);  
  
FD_CLR(int fd, fd_set *fds);    /* clear the bit for fd in fds */  
FD_ISSET(int fd, fd_set *fds); /* is the bit for fd in fds? */  
FD_SET(int fd, fd_set *fds);    /* turn on the bit for fd in fds */  
FD_ZERO(fd_set *fds);          /* clear all bits in fds */
```

- **maxfds**: number of descriptors to be tested
 - descriptors (0, 1, ... maxfds-1) will be tested
- **readfds**: a set of *fds* we want to check if data is available
 - returns a set of *fds* ready to read
 - if input argument is *NULL*, not interested in that condition
- **writelfds**: returns a set of *fds* ready to write
- **exceptfds**: returns a set of *fds* with exception conditions

Socket I/O: select()

```
int select(int maxfds, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);

struct timeval {
    long tv_sec;           /* seconds /
    long tv_usec;          /* microseconds */
}
```

- ***timeout***
 - if NULL, wait forever and return only when one of the descriptors is ready for I/O
 - otherwise, wait up to a fixed amount of time specified by *timeout*
 - if we don't want to wait at all, create a timeout structure with timer value equal to 0
- Refer to the man page for more information

Socket I/O: select()


- *select* allows synchronous I/O multiplexing

```
int s1, s2;                /* socket descriptors */
fd_set readfds;            /* used by select() */
int not_done = 1;

/* create and bind s1 and s2 */
while(not_done) {
    FD_ZERO(&readfds);      /* initialize the fd set */
    /*
    FD_SET(s1, &readfds);   /* add s1 to the fd set */
    FD_SET(s2, &readfds);   /* add s2 to the fd set */

    if(select(s2+1, &readfds, 0, 0, 0) < 0) {
        perror("select");
        exit(1);
    }
    if(FD_ISSET(s1, &readfds)) {
        recvfrom(s1, buf, sizeof(buf), ...);
        /* process buf */
    }
    /* do the same for s2 */
}
```


More Details About a Web Server



How can a web server manage multiple connections simultaneously?

Socket I/O: select()

```
int fd, next=0;                /* original socket */
int newfd[10];                 /* new socket descriptors */
int not_done = 1;

while(not_done) {
    fd_set readfds;
    FD_ZERO(&readfds);
    FD_SET(fd, &readfds);

    /* Now use FD_SET to initialize other newfd's
       that have already been returned by accept() */

    select(maxfd+1, &readfds, 0, 0, 0);
    if(FD_ISSET(fd, &readfds)) {
        newfd[next++] = accept(fd, ...);
    }
    /* do the following for each descriptor newfd[n] */
    if(FD_ISSET(newfd[n], &readfds)) {
        read(newfd[n], buf, sizeof(buf));
        /* process data */
    }
}
```

- **Now the web server can support multiple connections...**

Other useful functions

- `bzero(char* c, int n)`: 0's n bytes starting at c
- `gethostname(char *name, int len)`: gets the name of the current host
- `gethostbyaddr(char *addr, int len, int type)`: converts IP hostname to structure containing long integer
- `inet_addr(const char *cp)`: converts dotted-decimal char-string to long integer
- `inet_ntoa(const struct in_addr in)`: converts long to dotted-decimal notation
- Warning: check function assumptions about byte-ordering (host or network). Often, they assume parameters / return solutions in network byte-order